

q3

April 3, 2020

0.1 Question 3 - MNIST Classification using PyTorch

```
[0]: import torch
import numpy as np
```

Load Data Sets

```
[2]: from torchvision import datasets
import torchvision.transforms as transforms

# number of subprocesses to use for data loading
num_workers = 0
# how many samples per batch to load
batch_size = 20

# convert data to torch.FloatTensor
transform = transforms.ToTensor()

# choose the training and test datasets
train_data = datasets.MNIST(root='data', train=True,
                             download=True, transform=transform)
test_data = datasets.MNIST(root='data', train=False,
                             download=True, transform=transform)

# prepare data loaders
train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                             num_workers=num_workers)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,
                                             num_workers=num_workers)
```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to
data/MNIST/raw/train-images-idx3-ubyte.gz

HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))

Extracting data/MNIST/raw/train-images-idx3-ubyte.gz to data/MNIST/raw
Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to
data/MNIST/raw/train-labels-idx1-ubyte.gz

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

Extracting data/MNIST/raw/train-labels-idx1-ubyte.gz to data/MNIST/raw
Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to
data/MNIST/raw/t10k-images-idx3-ubyte.gz

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

Extracting data/MNIST/raw/t10k-images-idx3-ubyte.gz to data/MNIST/raw
Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to
data/MNIST/raw/t10k-labels-idx1-ubyte.gz

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```


Extracting data/MNIST/raw/t10k-labels-idx1-ubyte.gz to data/MNIST/raw
Processing...
Done!

Visualize a batch of Training Set

```
[3]: import matplotlib.pyplot as plt
      %matplotlib inline

      # obtain one batch of training images
      dataiter = iter(train_loader)
      images, labels = dataiter.next()
      images = images.numpy()

      # plot the images in the batch, along with the corresponding labels
      fig = plt.figure(figsize=(25, 4))
      for idx in np.arange(20):
          ax = fig.add_subplot(2, 20/2, idx+1, xticks=[], yticks=[])
          ax.imshow(np.squeeze(images[idx]), cmap='gray')
          # print out the correct label for each image
          # .item() gets the value contained in a Tensor
          ax.set_title(str(labels[idx].item()))
```



Define the Network Architecture

```
[0]: import torch.nn as nn
import torch.nn.functional as F

## Define the NN architecture
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(28 * 28, 512)
        # linear layer (n_hidden -> hidden_2)
        self.fc2 = nn.Linear(512, 512)
        # linear layer (n_hidden -> 10)
        self.fc3 = nn.Linear(512, 10)
        # dropout layer (p=0.2)
        # dropout prevents overfitting of data
        self.dropout = nn.Dropout(0.2)

    def forward(self, x):
        # flatten image input
        x = x.view(-1, 28 * 28)
        # add hidden layer, with relu activation function
        x = F.relu(self.fc1(x))
        return x
```

```
[5]: # initialize the NN
model = Net()
print(model)
```

```
Net(
  (fc1): Linear(in_features=784, out_features=512, bias=True)
  (fc2): Linear(in_features=512, out_features=512, bias=True)
  (fc3): Linear(in_features=512, out_features=10, bias=True)
  (dropout): Dropout(p=0.2, inplace=False)
)
```

Specify Loss Function and Optimizer

```
[0]: criterion = nn.CrossEntropyLoss()

# specify optimizer
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
```

Train the Network

```

[7]: n_epochs = 30 # suggest training between 20-50 epochs

model.train() # prep model for training

for epoch in range(n_epochs):
    # monitor training loss
    train_loss = 0.0

    #####
    # train the model #
    #####
    for data, target in train_loader:
        # clear the gradients of all optimized variables
        optimizer.zero_grad()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # backward pass: compute gradient of the loss with respect to model_
        ↪ parameters
        loss.backward()
        # perform a single optimization step (parameter update)
        optimizer.step()
        # update running training loss
        train_loss += loss.item()*data.size(0)

    # print training statistics
    # calculate average loss over an epoch
    train_loss = train_loss/len(train_loader.dataset)

    print('Epoch: {} \tTraining Loss: {:.6f}'.format(
        epoch+1,
        train_loss
    ))

```

```

Epoch: 1      Training Loss: 0.804222
Epoch: 2      Training Loss: 0.411707
Epoch: 3      Training Loss: 0.370065
Epoch: 4      Training Loss: 0.348959
Epoch: 5      Training Loss: 0.335502
Epoch: 6      Training Loss: 0.325918
Epoch: 7      Training Loss: 0.318626
Epoch: 8      Training Loss: 0.312826
Epoch: 9      Training Loss: 0.308064
Epoch: 10     Training Loss: 0.304058
Epoch: 11     Training Loss: 0.300624
Epoch: 12     Training Loss: 0.297635

```

Epoch: 13	Training Loss: 0.294999
Epoch: 14	Training Loss: 0.292651
Epoch: 15	Training Loss: 0.290539
Epoch: 16	Training Loss: 0.288626
Epoch: 17	Training Loss: 0.286881
Epoch: 18	Training Loss: 0.285279
Epoch: 19	Training Loss: 0.283802
Epoch: 20	Training Loss: 0.282434
Epoch: 21	Training Loss: 0.281160
Epoch: 22	Training Loss: 0.279970
Epoch: 23	Training Loss: 0.278856
Epoch: 24	Training Loss: 0.277808
Epoch: 25	Training Loss: 0.276820
Epoch: 26	Training Loss: 0.275887
Epoch: 27	Training Loss: 0.275002
Epoch: 28	Training Loss: 0.274163
Epoch: 29	Training Loss: 0.273365
Epoch: 30	Training Loss: 0.272604

Test the Trained Network

```
[8]: test_loss = 0.0
class_correct = list(0. for i in range(10))
class_total = list(0. for i in range(10))

model.eval() # prep model for *evaluation*

for data, target in test_loader:
    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data)
    # calculate the loss
    loss = criterion(output, target)
    # update test loss
    test_loss += loss.item()*data.size(0)
    # convert output probabilities to predicted class
    _, pred = torch.max(output, 1)
    # compare predictions to true label
    correct = np.squeeze(pred.eq(target.data.view_as(pred)))
    # calculate test accuracy for each object class
    for i in range(batch_size):
        label = target.data[i]
        class_correct[label] += correct[i].item()
        class_total[label] += 1

# calculate and print avg test loss
test_loss = test_loss/len(test_loader.dataset)
print('Test Loss: {:.6f}\n'.format(test_loss))
```

```

for i in range(10):
    if class_total[i] > 0:
        print('Test Accuracy of %5s: %2d%% (%2d/%2d)' % (
            str(i), 100 * class_correct[i] / class_total[i],
            np.sum(class_correct[i]), np.sum(class_total[i])))
    else:
        print('Test Accuracy of %5s: N/A (no training examples)' % (classes[i]))

print('\nTest Accuracy (Overall): %2d%% (%2d/%2d)' % (
    100. * np.sum(class_correct) / np.sum(class_total),
    np.sum(class_correct), np.sum(class_total)))

```

Test Loss: 0.274007

```

Test Accuracy of    0: 98% (962/980)
Test Accuracy of    1: 97% (1109/1135)
Test Accuracy of    2: 88% (918/1032)
Test Accuracy of    3: 90% (914/1010)
Test Accuracy of    4: 92% (909/982)
Test Accuracy of    5: 87% (781/892)
Test Accuracy of    6: 94% (910/958)
Test Accuracy of    7: 92% (946/1028)
Test Accuracy of    8: 88% (864/974)
Test Accuracy of    9: 90% (916/1009)

```

Test Accuracy (Overall): 92% (9229/10000)

Visualize Sample Test Results


```

[9]: dataiter = iter(test_loader)
    images, labels = dataiter.next()

    # get sample outputs
    output = model(images)
    # convert output probabilities to predicted class
    _, preds = torch.max(output, 1)
    # prep images for display
    images = images.numpy()

    # plot the images in the batch, along with predicted and true labels
    fig = plt.figure(figsize=(25, 4))
    for idx in np.arange(20):
        ax = fig.add_subplot(2, 20/2, idx+1, xticks=[], yticks=[])
        ax.imshow(np.squeeze(images[idx]), cmap='gray')
        ax.set_title("{} ({}).format(str(preds[idx].item()), str(labels[idx].
→item()))),
                                color=("green" if preds[idx]==labels[idx] else "red"))

```



0.2 MNIST Classification using CNN

```
[0]: import numpy as np # to handle matrix and data operation
import pandas as pd # to read csv and handle dataframe

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.utils.data
from torch.autograd import Variable

from sklearn.model_selection import train_test_split
```

```
[0]: from keras.datasets import mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
[0]: x_train = x_train.reshape(60000,784)
x_test = x_test.reshape(10000,784)
```

```
[0]: BATCH_SIZE = 32
```

```
[0]: torch_X_train = torch.from_numpy(x_train).type(torch.LongTensor)
torch_y_train = torch.from_numpy(y_train).type(torch.LongTensor)
```

```
[0]: torch_X_test = torch.from_numpy(x_test).type(torch.LongTensor)
torch_y_test = torch.from_numpy(y_test).type(torch.LongTensor)
```

```
[0]: train = torch.utils.data.TensorDataset(torch_X_train,torch_y_train)
test = torch.utils.data.TensorDataset(torch_X_test,torch_y_test)
```

```
[0]: train_loader = torch.utils.data.DataLoader(train, batch_size = BATCH_SIZE,
    ↪shuffle = False)
test_loader = torch.utils.data.DataLoader(test, batch_size = BATCH_SIZE,
    ↪shuffle = False)
```

```
[59]: class MLP(nn.Module):
    def __init__(self):
        super(MLP, self).__init__()
        self.linear1 = nn.Linear(784,250)
        self.linear2 = nn.Linear(250,100)
        self.linear3 = nn.Linear(100,10)

    def forward(self,X):
        X = F.relu(self.linear1(X))
        X = F.relu(self.linear2(X))
        X = self.linear3(X)
        return F.log_softmax(X, dim=1)

mlp = MLP()
print(mlp)
```

```
MLP(
  (linear1): Linear(in_features=784, out_features=250, bias=True)
  (linear2): Linear(in_features=250, out_features=100, bias=True)
  (linear3): Linear(in_features=100, out_features=10, bias=True)
)
```

```
[0]: def fit(model, train_loader):
    optimizer = torch.optim.Adam(model.parameters())#,lr=0.001, betas=(0.9,0.
    ↪999))
    error = nn.CrossEntropyLoss()
    EPOCHS = 5
    model.train()
    for epoch in range(EPOCHS):
        correct = 0
        for batch_idx, (X_batch, y_batch) in enumerate(train_loader):
            var_X_batch = Variable(X_batch).float()
            var_y_batch = Variable(y_batch)
            optimizer.zero_grad()
            output = model(var_X_batch)
            loss = error(output, var_y_batch)
            loss.backward()
            optimizer.step()

            # Total correct predictions
            predicted = torch.max(output.data, 1)[1]
            correct += (predicted == var_y_batch).sum()
            #print(correct)
            if batch_idx % 50 == 0:
                print('Epoch : {} [{} / {}] {:.0f}%]\tLoss: {:.6f}\t Accuracy:{:.
                ↪3f}%'.format(
```



```

epoch, batch_idx*len(X_batch), len(train_loader.dataset),
↪100.*batch_idx / len(train_loader), (loss.data), float(correct*100) /
↪float(BATCH_SIZE*(batch_idx+1)))

```

```
[67]: fit(mlp, train_loader)
```

Epoch : 0	[0/60000 (0%)]	Loss: 0.028779	Accuracy:100.000%
Epoch : 0	[1600/60000 (3%)]	Loss: 0.135546	Accuracy:96.875%
Epoch : 0	[3200/60000 (5%)]	Loss: 0.161492	Accuracy:97.061%
Epoch : 0	[4800/60000 (8%)]	Loss: 0.029837	Accuracy:97.144%
Epoch : 0	[6400/60000 (11%)]	Loss: 0.251311	Accuracy:97.233%
Epoch : 0	[8000/60000 (13%)]	Loss: 0.288707	Accuracy:96.987%
Epoch : 0	[9600/60000 (16%)]	Loss: 0.001341	Accuracy:96.564%
Epoch : 0	[11200/60000 (19%)]	Loss: 0.147621	Accuracy:96.270%
Epoch : 0	[12800/60000 (21%)]	Loss: 0.162637	Accuracy:96.213%
Epoch : 0	[14400/60000 (24%)]	Loss: 0.040754	Accuracy:96.182%
Epoch : 0	[16000/60000 (27%)]	Loss: 0.053642	Accuracy:96.214%
Epoch : 0	[17600/60000 (29%)]	Loss: 0.021809	Accuracy:96.257%
Epoch : 0	[19200/60000 (32%)]	Loss: 0.007951	Accuracy:96.339%
Epoch : 0	[20800/60000 (35%)]	Loss: 0.117186	Accuracy:96.376%
Epoch : 0	[22400/60000 (37%)]	Loss: 0.001032	Accuracy:96.438%
Epoch : 0	[24000/60000 (40%)]	Loss: 0.240103	Accuracy:96.451%
Epoch : 0	[25600/60000 (43%)]	Loss: 0.097887	Accuracy:96.497%
Epoch : 0	[27200/60000 (45%)]	Loss: 0.023727	Accuracy:96.500%
Epoch : 0	[28800/60000 (48%)]	Loss: 0.015119	Accuracy:96.528%
Epoch : 0	[30400/60000 (51%)]	Loss: 0.070183	Accuracy:96.537%
Epoch : 0	[32000/60000 (53%)]	Loss: 0.012500	Accuracy:96.544%
Epoch : 0	[33600/60000 (56%)]	Loss: 0.110607	Accuracy:96.530%
Epoch : 0	[35200/60000 (59%)]	Loss: 0.357193	Accuracy:96.540%
Epoch : 0	[36800/60000 (61%)]	Loss: 0.010914	Accuracy:96.549%
Epoch : 0	[38400/60000 (64%)]	Loss: 0.176304	Accuracy:96.555%
Epoch : 0	[40000/60000 (67%)]	Loss: 0.098055	Accuracy:96.530%
Epoch : 0	[41600/60000 (69%)]	Loss: 0.044533	Accuracy:96.539%
Epoch : 0	[43200/60000 (72%)]	Loss: 0.099969	Accuracy:96.551%
Epoch : 0	[44800/60000 (75%)]	Loss: 0.024746	Accuracy:96.576%
Epoch : 0	[46400/60000 (77%)]	Loss: 0.333618	Accuracy:96.550%
Epoch : 0	[48000/60000 (80%)]	Loss: 0.118122	Accuracy:96.565%
Epoch : 0	[49600/60000 (83%)]	Loss: 0.100130	Accuracy:96.555%
Epoch : 0	[51200/60000 (85%)]	Loss: 0.101556	Accuracy:96.549%
Epoch : 0	[52800/60000 (88%)]	Loss: 0.163493	Accuracy:96.555%
Epoch : 0	[54400/60000 (91%)]	Loss: 0.001940	Accuracy:96.550%
Epoch : 0	[56000/60000 (93%)]	Loss: 0.151699	Accuracy:96.552%
Epoch : 0	[57600/60000 (96%)]	Loss: 0.085508	Accuracy:96.577%
Epoch : 0	[59200/60000 (99%)]	Loss: 0.006413	Accuracy:96.607%
Epoch : 1	[0/60000 (0%)]	Loss: 0.066838	Accuracy:96.875%
Epoch : 1	[1600/60000 (3%)]	Loss: 0.038393	Accuracy:95.895%
Epoch : 1	[3200/60000 (5%)]	Loss: 0.088629	Accuracy:96.566%

Epoch : 1	[4800/60000 (8%)]	Loss: 0.073845	Accuracy:96.978%
Epoch : 1	[6400/60000 (11%)]	Loss: 0.323887	Accuracy:97.015%
Epoch : 1	[8000/60000 (13%)]	Loss: 0.080831	Accuracy:97.049%
Epoch : 1	[9600/60000 (16%)]	Loss: 0.028484	Accuracy:97.103%
Epoch : 1	[11200/60000 (19%)]	Loss: 0.217653	Accuracy:97.044%
Epoch : 1	[12800/60000 (21%)]	Loss: 0.015726	Accuracy:97.015%
Epoch : 1	[14400/60000 (24%)]	Loss: 0.031854	Accuracy:96.993%
Epoch : 1	[16000/60000 (27%)]	Loss: 0.102694	Accuracy:97.050%
Epoch : 1	[17600/60000 (29%)]	Loss: 0.314560	Accuracy:97.164%
Epoch : 1	[19200/60000 (32%)]	Loss: 0.043562	Accuracy:97.192%
Epoch : 1	[20800/60000 (35%)]	Loss: 0.078116	Accuracy:97.211%
Epoch : 1	[22400/60000 (37%)]	Loss: 0.010395	Accuracy:97.214%
Epoch : 1	[24000/60000 (40%)]	Loss: 0.008545	Accuracy:97.241%
Epoch : 1	[25600/60000 (43%)]	Loss: 0.087680	Accuracy:97.285%
Epoch : 1	[27200/60000 (45%)]	Loss: 0.016307	Accuracy:97.264%
Epoch : 1	[28800/60000 (48%)]	Loss: 0.070700	Accuracy:97.295%
Epoch : 1	[30400/60000 (51%)]	Loss: 0.283281	Accuracy:97.315%
Epoch : 1	[32000/60000 (53%)]	Loss: 0.124514	Accuracy:97.331%
Epoch : 1	[33600/60000 (56%)]	Loss: 0.077984	Accuracy:97.265%
Epoch : 1	[35200/60000 (59%)]	Loss: 0.018130	Accuracy:97.272%
Epoch : 1	[36800/60000 (61%)]	Loss: 0.063739	Accuracy:97.277%
Epoch : 1	[38400/60000 (64%)]	Loss: 0.262766	Accuracy:97.250%
Epoch : 1	[40000/60000 (67%)]	Loss: 0.040512	Accuracy:97.252%
Epoch : 1	[41600/60000 (69%)]	Loss: 0.088825	Accuracy:97.235%
Epoch : 1	[43200/60000 (72%)]	Loss: 0.161225	Accuracy:97.213%
Epoch : 1	[44800/60000 (75%)]	Loss: 0.050391	Accuracy:97.232%
Epoch : 1	[46400/60000 (77%)]	Loss: 0.577136	Accuracy:97.194%
Epoch : 1	[48000/60000 (80%)]	Loss: 0.154128	Accuracy:97.183%
Epoch : 1	[49600/60000 (83%)]	Loss: 0.199059	Accuracy:97.155%
Epoch : 1	[51200/60000 (85%)]	Loss: 0.163394	Accuracy:97.146%
Epoch : 1	[52800/60000 (88%)]	Loss: 0.163188	Accuracy:97.161%
Epoch : 1	[54400/60000 (91%)]	Loss: 0.007887	Accuracy:97.163%
Epoch : 1	[56000/60000 (93%)]	Loss: 0.123494	Accuracy:97.150%
Epoch : 1	[57600/60000 (96%)]	Loss: 0.084630	Accuracy:97.167%
Epoch : 1	[59200/60000 (99%)]	Loss: 0.041318	Accuracy:97.197%
Epoch : 2	[0/60000 (0%)]	Loss: 0.101924	Accuracy:93.750%
Epoch : 2	[1600/60000 (3%)]	Loss: 0.192111	Accuracy:96.078%
Epoch : 2	[3200/60000 (5%)]	Loss: 0.070167	Accuracy:97.123%
Epoch : 2	[4800/60000 (8%)]	Loss: 0.023145	Accuracy:97.413%
Epoch : 2	[6400/60000 (11%)]	Loss: 0.164375	Accuracy:97.217%
Epoch : 2	[8000/60000 (13%)]	Loss: 0.206731	Accuracy:97.261%
Epoch : 2	[9600/60000 (16%)]	Loss: 0.013814	Accuracy:97.228%
Epoch : 2	[11200/60000 (19%)]	Loss: 0.189980	Accuracy:97.302%
Epoch : 2	[12800/60000 (21%)]	Loss: 0.094958	Accuracy:97.319%
Epoch : 2	[14400/60000 (24%)]	Loss: 0.015213	Accuracy:97.367%
Epoch : 2	[16000/60000 (27%)]	Loss: 0.093988	Accuracy:97.287%
Epoch : 2	[17600/60000 (29%)]	Loss: 0.621408	Accuracy:97.221%
Epoch : 2	[19200/60000 (32%)]	Loss: 0.108640	Accuracy:97.223%

Epoch : 2	[20800/60000 (35%)]	Loss: 0.021036	Accuracy:97.216%
Epoch : 2	[22400/60000 (37%)]	Loss: 0.001247	Accuracy:97.209%
Epoch : 2	[24000/60000 (40%)]	Loss: 0.015268	Accuracy:97.233%
Epoch : 2	[25600/60000 (43%)]	Loss: 0.101386	Accuracy:97.261%
Epoch : 2	[27200/60000 (45%)]	Loss: 0.131324	Accuracy:97.239%
Epoch : 2	[28800/60000 (48%)]	Loss: 0.069168	Accuracy:97.243%
Epoch : 2	[30400/60000 (51%)]	Loss: 0.025075	Accuracy:97.282%
Epoch : 2	[32000/60000 (53%)]	Loss: 0.016453	Accuracy:97.300%
Epoch : 2	[33600/60000 (56%)]	Loss: 0.020328	Accuracy:97.273%
Epoch : 2	[35200/60000 (59%)]	Loss: 0.872610	Accuracy:97.298%
Epoch : 2	[36800/60000 (61%)]	Loss: 0.101088	Accuracy:97.296%
Epoch : 2	[38400/60000 (64%)]	Loss: 0.088246	Accuracy:97.294%
Epoch : 2	[40000/60000 (67%)]	Loss: 0.119826	Accuracy:97.302%
Epoch : 2	[41600/60000 (69%)]	Loss: 0.052212	Accuracy:97.324%
Epoch : 2	[43200/60000 (72%)]	Loss: 0.325394	Accuracy:97.305%
Epoch : 2	[44800/60000 (75%)]	Loss: 0.062563	Accuracy:97.310%
Epoch : 2	[46400/60000 (77%)]	Loss: 0.133759	Accuracy:97.273%
Epoch : 2	[48000/60000 (80%)]	Loss: 0.219807	Accuracy:97.271%
Epoch : 2	[49600/60000 (83%)]	Loss: 0.019319	Accuracy:97.288%
Epoch : 2	[51200/60000 (85%)]	Loss: 0.045573	Accuracy:97.273%
Epoch : 2	[52800/60000 (88%)]	Loss: 0.129453	Accuracy:97.265%
Epoch : 2	[54400/60000 (91%)]	Loss: 0.037412	Accuracy:97.270%
Epoch : 2	[56000/60000 (93%)]	Loss: 0.061451	Accuracy:97.275%
Epoch : 2	[57600/60000 (96%)]	Loss: 0.268131	Accuracy:97.278%
Epoch : 2	[59200/60000 (99%)]	Loss: 0.029599	Accuracy:97.316%
Epoch : 3	[0/60000 (0%)]	Loss: 0.151179	Accuracy:93.750%
Epoch : 3	[1600/60000 (3%)]	Loss: 0.128003	Accuracy:96.446%
Epoch : 3	[3200/60000 (5%)]	Loss: 0.006657	Accuracy:97.308%
Epoch : 3	[4800/60000 (8%)]	Loss: 0.015271	Accuracy:97.392%
Epoch : 3	[6400/60000 (11%)]	Loss: 0.127014	Accuracy:97.404%
Epoch : 3	[8000/60000 (13%)]	Loss: 0.137492	Accuracy:97.323%
Epoch : 3	[9600/60000 (16%)]	Loss: 0.011746	Accuracy:97.290%
Epoch : 3	[11200/60000 (19%)]	Loss: 0.193034	Accuracy:97.391%
Epoch : 3	[12800/60000 (21%)]	Loss: 0.151988	Accuracy:97.506%
Epoch : 3	[14400/60000 (24%)]	Loss: 0.001826	Accuracy:97.561%
Epoch : 3	[16000/60000 (27%)]	Loss: 0.038850	Accuracy:97.493%
Epoch : 3	[17600/60000 (29%)]	Loss: 0.005036	Accuracy:97.533%
Epoch : 3	[19200/60000 (32%)]	Loss: 0.041748	Accuracy:97.577%
Epoch : 3	[20800/60000 (35%)]	Loss: 0.061989	Accuracy:97.552%
Epoch : 3	[22400/60000 (37%)]	Loss: 0.000083	Accuracy:97.562%
Epoch : 3	[24000/60000 (40%)]	Loss: 0.010684	Accuracy:97.562%
Epoch : 3	[25600/60000 (43%)]	Loss: 0.035633	Accuracy:97.554%
Epoch : 3	[27200/60000 (45%)]	Loss: 0.003687	Accuracy:97.536%
Epoch : 3	[28800/60000 (48%)]	Loss: 0.059841	Accuracy:97.562%
Epoch : 3	[30400/60000 (51%)]	Loss: 0.055981	Accuracy:97.581%
Epoch : 3	[32000/60000 (53%)]	Loss: 0.057697	Accuracy:97.596%
Epoch : 3	[33600/60000 (56%)]	Loss: 0.099345	Accuracy:97.589%
Epoch : 3	[35200/60000 (59%)]	Loss: 0.004503	Accuracy:97.593%

Epoch : 3	[36800/60000 (61%)]	Loss: 0.132729	Accuracy:97.597%
Epoch : 3	[38400/60000 (64%)]	Loss: 0.078958	Accuracy:97.580%
Epoch : 3	[40000/60000 (67%)]	Loss: 0.049452	Accuracy:97.577%
Epoch : 3	[41600/60000 (69%)]	Loss: 0.001719	Accuracy:97.574%
Epoch : 3	[43200/60000 (72%)]	Loss: 0.175857	Accuracy:97.580%
Epoch : 3	[44800/60000 (75%)]	Loss: 0.034069	Accuracy:97.589%
Epoch : 3	[46400/60000 (77%)]	Loss: 0.021272	Accuracy:97.599%
Epoch : 3	[48000/60000 (80%)]	Loss: 0.064088	Accuracy:97.620%
Epoch : 3	[49600/60000 (83%)]	Loss: 0.003137	Accuracy:97.625%
Epoch : 3	[51200/60000 (85%)]	Loss: 0.014169	Accuracy:97.628%
Epoch : 3	[52800/60000 (88%)]	Loss: 0.221221	Accuracy:97.625%
Epoch : 3	[54400/60000 (91%)]	Loss: 0.001268	Accuracy:97.612%
Epoch : 3	[56000/60000 (93%)]	Loss: 0.063965	Accuracy:97.596%
Epoch : 3	[57600/60000 (96%)]	Loss: 0.254212	Accuracy:97.604%
Epoch : 3	[59200/60000 (99%)]	Loss: 0.000153	Accuracy:97.633%
Epoch : 4	[0/60000 (0%)]	Loss: 0.077204	Accuracy:96.875%
Epoch : 4	[1600/60000 (3%)]	Loss: 0.160978	Accuracy:97.059%
Epoch : 4	[3200/60000 (5%)]	Loss: 0.022128	Accuracy:97.463%
Epoch : 4	[4800/60000 (8%)]	Loss: 0.087849	Accuracy:97.641%
Epoch : 4	[6400/60000 (11%)]	Loss: 0.282174	Accuracy:97.512%
Epoch : 4	[8000/60000 (13%)]	Loss: 0.022924	Accuracy:97.572%
Epoch : 4	[9600/60000 (16%)]	Loss: 0.057789	Accuracy:97.456%
Epoch : 4	[11200/60000 (19%)]	Loss: 0.173873	Accuracy:97.480%
Epoch : 4	[12800/60000 (21%)]	Loss: 0.177891	Accuracy:97.498%
Epoch : 4	[14400/60000 (24%)]	Loss: 0.058999	Accuracy:97.568%
Epoch : 4	[16000/60000 (27%)]	Loss: 0.087360	Accuracy:97.493%
Epoch : 4	[17600/60000 (29%)]	Loss: 0.002993	Accuracy:97.544%
Epoch : 4	[19200/60000 (32%)]	Loss: 0.007460	Accuracy:97.577%
Epoch : 4	[20800/60000 (35%)]	Loss: 0.172938	Accuracy:97.576%
Epoch : 4	[22400/60000 (37%)]	Loss: 0.003647	Accuracy:97.597%
Epoch : 4	[24000/60000 (40%)]	Loss: 0.003558	Accuracy:97.616%
Epoch : 4	[25600/60000 (43%)]	Loss: 0.123521	Accuracy:97.651%
Epoch : 4	[27200/60000 (45%)]	Loss: 0.062108	Accuracy:97.650%
Epoch : 4	[28800/60000 (48%)]	Loss: 0.079932	Accuracy:97.683%
Epoch : 4	[30400/60000 (51%)]	Loss: 0.002716	Accuracy:97.703%
Epoch : 4	[32000/60000 (53%)]	Loss: 0.218044	Accuracy:97.687%
Epoch : 4	[33600/60000 (56%)]	Loss: 0.009926	Accuracy:97.693%
Epoch : 4	[35200/60000 (59%)]	Loss: 0.001960	Accuracy:97.707%
Epoch : 4	[36800/60000 (61%)]	Loss: 0.003491	Accuracy:97.706%
Epoch : 4	[38400/60000 (64%)]	Loss: 0.193213	Accuracy:97.684%
Epoch : 4	[40000/60000 (67%)]	Loss: 0.121122	Accuracy:97.677%
Epoch : 4	[41600/60000 (69%)]	Loss: 0.058635	Accuracy:97.672%
Epoch : 4	[43200/60000 (72%)]	Loss: 0.121910	Accuracy:97.668%
Epoch : 4	[44800/60000 (75%)]	Loss: 0.020887	Accuracy:97.711%
Epoch : 4	[46400/60000 (77%)]	Loss: 0.099990	Accuracy:97.676%
Epoch : 4	[48000/60000 (80%)]	Loss: 0.167490	Accuracy:97.677%
Epoch : 4	[49600/60000 (83%)]	Loss: 0.096079	Accuracy:97.671%
Epoch : 4	[51200/60000 (85%)]	Loss: 0.009618	Accuracy:97.687%

```
Epoch : 4 [52800/60000 (88%)]    Loss: 0.322359    Accuracy:97.714%
Epoch : 4 [54400/60000 (91%)]    Loss: 0.001315    Accuracy:97.702%
Epoch : 4 [56000/60000 (93%)]    Loss: 0.065532    Accuracy:97.717%
Epoch : 4 [57600/60000 (96%)]    Loss: 0.120293    Accuracy:97.729%
Epoch : 4 [59200/60000 (99%)]    Loss: 0.000217    Accuracy:97.758%
```

```
[68]: def evaluate(model):
    #model = mlp
    correct = 0
    for test_imgs, test_labels in test_loader:
        #print(test_imgs.shape)
        test_imgs = Variable(test_imgs).float()
        output = model(test_imgs)
        predicted = torch.max(output,1)[1]
        correct += (predicted == test_labels).sum()
    print("Test accuracy:{:.3f}% ".format( float(correct) /
    ↪(len(test_loader)*BATCH_SIZE)))
    evaluate(mlp)
```

Test accuracy:0.971%

##Since a CNN needs a image shape as input let's reshape our flatten images to real image

```
[69]: torch_X_train = torch_X_train.view(-1, 1,28,28).float()
    torch_X_test = torch_X_test.view(-1,1,28,28).float()
    print(torch_X_train.shape)
    print(torch_X_test.shape)

    # Pytorch train and test sets
    train = torch.utils.data.TensorDataset(torch_X_train,torch_y_train)
    test = torch.utils.data.TensorDataset(torch_X_test,torch_y_test)

    # data loader
    train_loader = torch.utils.data.DataLoader(train, batch_size = BATCH_SIZE,
    ↪shuffle = False)
    test_loader = torch.utils.data.DataLoader(test, batch_size = BATCH_SIZE,
    ↪shuffle = False)MB
```

torch.Size([60000, 1, 28, 28])

torch.Size([10000, 1, 28, 28])

```
[70]: class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=5)
        self.conv2 = nn.Conv2d(32, 32, kernel_size=5)
        self.conv3 = nn.Conv2d(32,64, kernel_size=5)
        self.fc1 = nn.Linear(3*3*64, 256)
```

```

        self.fc2 = nn.Linear(256, 10)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        #x = F.dropout(x, p=0.5, training=self.training)
        x = F.relu(F.max_pool2d(self.conv2(x), 2))
        x = F.dropout(x, p=0.5, training=self.training)
        x = F.relu(F.max_pool2d(self.conv3(x), 2))
        x = F.dropout(x, p=0.5, training=self.training)
        x = x.view(-1, 3*3*64)
        x = F.relu(self.fc1(x))
        x = F.dropout(x, training=self.training)
        x = self.fc2(x)
        return F.log_softmax(x, dim=1)

cnn = CNN()
print(cnn)

it = iter(train_loader)
X_batch, y_batch = next(it)
print(cnn.forward(X_batch).shape)

```

```

CNN(
  (conv1): Conv2d(1, 32, kernel_size=(5, 5), stride=(1, 1))
  (conv2): Conv2d(32, 32, kernel_size=(5, 5), stride=(1, 1))
  (conv3): Conv2d(32, 64, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear(in_features=576, out_features=256, bias=True)
  (fc2): Linear(in_features=256, out_features=10, bias=True)
)
torch.Size([32, 10])

```

```
[71]: fit(cnn, train_loader)
```

Epoch : 0 [0/60000 (0%)]	Loss: 24.703955	Accuracy: 6.250%
Epoch : 0 [1600/60000 (3%)]	Loss: 1.980315	Accuracy: 17.279%
Epoch : 0 [3200/60000 (5%)]	Loss: 1.460487	Accuracy: 29.115%
Epoch : 0 [4800/60000 (8%)]	Loss: 0.748759	Accuracy: 38.949%
Epoch : 0 [6400/60000 (11%)]	Loss: 1.238455	Accuracy: 45.553%
Epoch : 0 [8000/60000 (13%)]	Loss: 0.865389	Accuracy: 50.461%
Epoch : 0 [9600/60000 (16%)]	Loss: 0.764462	Accuracy: 54.485%
Epoch : 0 [11200/60000 (19%)]	Loss: 0.602340	Accuracy: 58.191%
Epoch : 0 [12800/60000 (21%)]	Loss: 0.548422	Accuracy: 60.793%
Epoch : 0 [14400/60000 (24%)]	Loss: 0.491586	Accuracy: 63.089%
Epoch : 0 [16000/60000 (27%)]	Loss: 0.702923	Accuracy: 64.814%
Epoch : 0 [17600/60000 (29%)]	Loss: 0.591373	Accuracy: 66.640%
Epoch : 0 [19200/60000 (32%)]	Loss: 0.427941	Accuracy: 68.225%
Epoch : 0 [20800/60000 (35%)]	Loss: 0.508253	Accuracy: 69.772%
Epoch : 0 [22400/60000 (37%)]	Loss: 0.220711	Accuracy: 71.122%

Epoch : 0	[24000/60000 (40%)]	Loss: 0.182511	Accuracy:72.121%
Epoch : 0	[25600/60000 (43%)]	Loss: 0.578731	Accuracy:73.174%
Epoch : 0	[27200/60000 (45%)]	Loss: 0.405716	Accuracy:74.104%
Epoch : 0	[28800/60000 (48%)]	Loss: 0.335528	Accuracy:75.035%
Epoch : 0	[30400/60000 (51%)]	Loss: 0.524058	Accuracy:75.710%
Epoch : 0	[32000/60000 (53%)]	Loss: 0.522807	Accuracy:76.346%
Epoch : 0	[33600/60000 (56%)]	Loss: 0.353773	Accuracy:77.019%
Epoch : 0	[35200/60000 (59%)]	Loss: 0.348661	Accuracy:77.677%
Epoch : 0	[36800/60000 (61%)]	Loss: 0.335008	Accuracy:78.307%
Epoch : 0	[38400/60000 (64%)]	Loss: 0.295484	Accuracy:78.841%
Epoch : 0	[40000/60000 (67%)]	Loss: 0.330925	Accuracy:79.359%
Epoch : 0	[41600/60000 (69%)]	Loss: 0.761555	Accuracy:79.790%
Epoch : 0	[43200/60000 (72%)]	Loss: 0.308168	Accuracy:80.204%
Epoch : 0	[44800/60000 (75%)]	Loss: 0.039641	Accuracy:80.661%
Epoch : 0	[46400/60000 (77%)]	Loss: 0.347005	Accuracy:81.032%
Epoch : 0	[48000/60000 (80%)]	Loss: 0.695292	Accuracy:81.404%
Epoch : 0	[49600/60000 (83%)]	Loss: 0.664559	Accuracy:81.742%
Epoch : 0	[51200/60000 (85%)]	Loss: 0.267460	Accuracy:82.083%
Epoch : 0	[52800/60000 (88%)]	Loss: 0.457143	Accuracy:82.418%
Epoch : 0	[54400/60000 (91%)]	Loss: 0.123103	Accuracy:82.742%
Epoch : 0	[56000/60000 (93%)]	Loss: 0.171987	Accuracy:83.040%
Epoch : 0	[57600/60000 (96%)]	Loss: 0.293099	Accuracy:83.320%
Epoch : 0	[59200/60000 (99%)]	Loss: 0.139004	Accuracy:83.656%
Epoch : 1	[0/60000 (0%)]	Loss: 0.197437	Accuracy:93.750%
Epoch : 1	[1600/60000 (3%)]	Loss: 0.182057	Accuracy:93.137%
Epoch : 1	[3200/60000 (5%)]	Loss: 0.449420	Accuracy:94.028%
Epoch : 1	[4800/60000 (8%)]	Loss: 0.116307	Accuracy:93.729%
Epoch : 1	[6400/60000 (11%)]	Loss: 0.029951	Accuracy:93.797%
Epoch : 1	[8000/60000 (13%)]	Loss: 0.101050	Accuracy:93.875%
Epoch : 1	[9600/60000 (16%)]	Loss: 0.333615	Accuracy:93.625%
Epoch : 1	[11200/60000 (19%)]	Loss: 0.426801	Accuracy:93.643%
Epoch : 1	[12800/60000 (21%)]	Loss: 0.298322	Accuracy:93.462%
Epoch : 1	[14400/60000 (24%)]	Loss: 0.108020	Accuracy:93.528%
Epoch : 1	[16000/60000 (27%)]	Loss: 0.474114	Accuracy:93.469%
Epoch : 1	[17600/60000 (29%)]	Loss: 0.277469	Accuracy:93.455%
Epoch : 1	[19200/60000 (32%)]	Loss: 0.163615	Accuracy:93.506%
Epoch : 1	[20800/60000 (35%)]	Loss: 0.227395	Accuracy:93.568%
Epoch : 1	[22400/60000 (37%)]	Loss: 0.010190	Accuracy:93.630%
Epoch : 1	[24000/60000 (40%)]	Loss: 0.085058	Accuracy:93.633%
Epoch : 1	[25600/60000 (43%)]	Loss: 0.092341	Accuracy:93.606%
Epoch : 1	[27200/60000 (45%)]	Loss: 0.161767	Accuracy:93.548%
Epoch : 1	[28800/60000 (48%)]	Loss: 0.029576	Accuracy:93.573%
Epoch : 1	[30400/60000 (51%)]	Loss: 0.205116	Accuracy:93.589%
Epoch : 1	[32000/60000 (53%)]	Loss: 0.252389	Accuracy:93.581%
Epoch : 1	[33600/60000 (56%)]	Loss: 0.318719	Accuracy:93.583%
Epoch : 1	[35200/60000 (59%)]	Loss: 0.124187	Accuracy:93.583%
Epoch : 1	[36800/60000 (61%)]	Loss: 0.110199	Accuracy:93.636%
Epoch : 1	[38400/60000 (64%)]	Loss: 0.135114	Accuracy:93.633%

Epoch : 1	[40000/60000 (67%)]	Loss: 0.337480	Accuracy:93.608%
Epoch : 1	[41600/60000 (69%)]	Loss: 0.151904	Accuracy:93.637%
Epoch : 1	[43200/60000 (72%)]	Loss: 0.201784	Accuracy:93.651%
Epoch : 1	[44800/60000 (75%)]	Loss: 0.120788	Accuracy:93.670%
Epoch : 1	[46400/60000 (77%)]	Loss: 0.631979	Accuracy:93.638%
Epoch : 1	[48000/60000 (80%)]	Loss: 0.330198	Accuracy:93.644%
Epoch : 1	[49600/60000 (83%)]	Loss: 0.092436	Accuracy:93.667%
Epoch : 1	[51200/60000 (85%)]	Loss: 0.046334	Accuracy:93.693%
Epoch : 1	[52800/60000 (88%)]	Loss: 0.316572	Accuracy:93.701%
Epoch : 1	[54400/60000 (91%)]	Loss: 0.114109	Accuracy:93.695%
Epoch : 1	[56000/60000 (93%)]	Loss: 0.394695	Accuracy:93.727%
Epoch : 1	[57600/60000 (96%)]	Loss: 0.083526	Accuracy:93.792%
Epoch : 1	[59200/60000 (99%)]	Loss: 0.015199	Accuracy:93.860%
Epoch : 2	[0/60000 (0%)]	Loss: 0.217455	Accuracy:90.625%
Epoch : 2	[1600/60000 (3%)]	Loss: 0.406698	Accuracy:94.056%
Epoch : 2	[3200/60000 (5%)]	Loss: 0.423212	Accuracy:94.864%
Epoch : 2	[4800/60000 (8%)]	Loss: 0.065729	Accuracy:94.454%
Epoch : 2	[6400/60000 (11%)]	Loss: 0.006236	Accuracy:94.325%
Epoch : 2	[8000/60000 (13%)]	Loss: 0.204239	Accuracy:94.410%
Epoch : 2	[9600/60000 (16%)]	Loss: 0.303900	Accuracy:94.248%
Epoch : 2	[11200/60000 (19%)]	Loss: 0.255757	Accuracy:94.436%
Epoch : 2	[12800/60000 (21%)]	Loss: 0.152741	Accuracy:94.459%
Epoch : 2	[14400/60000 (24%)]	Loss: 0.082564	Accuracy:94.471%
Epoch : 2	[16000/60000 (27%)]	Loss: 0.283941	Accuracy:94.374%
Epoch : 2	[17600/60000 (29%)]	Loss: 0.014378	Accuracy:94.380%
Epoch : 2	[19200/60000 (32%)]	Loss: 0.083600	Accuracy:94.452%
Epoch : 2	[20800/60000 (35%)]	Loss: 0.360928	Accuracy:94.556%
Epoch : 2	[22400/60000 (37%)]	Loss: 0.102351	Accuracy:94.593%
Epoch : 2	[24000/60000 (40%)]	Loss: 0.109202	Accuracy:94.653%
Epoch : 2	[25600/60000 (43%)]	Loss: 0.092279	Accuracy:94.714%
Epoch : 2	[27200/60000 (45%)]	Loss: 0.241476	Accuracy:94.679%
Epoch : 2	[28800/60000 (48%)]	Loss: 0.107590	Accuracy:94.704%
Epoch : 2	[30400/60000 (51%)]	Loss: 0.188392	Accuracy:94.759%
Epoch : 2	[32000/60000 (53%)]	Loss: 0.298663	Accuracy:94.699%
Epoch : 2	[33600/60000 (56%)]	Loss: 0.165397	Accuracy:94.690%
Epoch : 2	[35200/60000 (59%)]	Loss: 0.109414	Accuracy:94.738%
Epoch : 2	[36800/60000 (61%)]	Loss: 0.111880	Accuracy:94.790%
Epoch : 2	[38400/60000 (64%)]	Loss: 0.247883	Accuracy:94.786%
Epoch : 2	[40000/60000 (67%)]	Loss: 0.241609	Accuracy:94.799%
Epoch : 2	[41600/60000 (69%)]	Loss: 0.274725	Accuracy:94.809%
Epoch : 2	[43200/60000 (72%)]	Loss: 0.335197	Accuracy:94.828%
Epoch : 2	[44800/60000 (75%)]	Loss: 0.005174	Accuracy:94.872%
Epoch : 2	[46400/60000 (77%)]	Loss: 1.024677	Accuracy:94.889%
Epoch : 2	[48000/60000 (80%)]	Loss: 0.283337	Accuracy:94.860%
Epoch : 2	[49600/60000 (83%)]	Loss: 0.264172	Accuracy:94.854%
Epoch : 2	[51200/60000 (85%)]	Loss: 0.145659	Accuracy:94.857%
Epoch : 2	[52800/60000 (88%)]	Loss: 0.662844	Accuracy:94.893%
Epoch : 2	[54400/60000 (91%)]	Loss: 0.017660	Accuracy:94.922%

Epoch : 2	[56000/60000 (93%)]	Loss: 0.118975	Accuracy:94.960%
Epoch : 2	[57600/60000 (96%)]	Loss: 0.390172	Accuracy:94.961%
Epoch : 2	[59200/60000 (99%)]	Loss: 0.000941	Accuracy:95.016%
Epoch : 3	[0/60000 (0%)]	Loss: 0.039304	Accuracy:96.875%
Epoch : 3	[1600/60000 (3%)]	Loss: 0.274808	Accuracy:94.424%
Epoch : 3	[3200/60000 (5%)]	Loss: 0.082404	Accuracy:95.142%
Epoch : 3	[4800/60000 (8%)]	Loss: 0.180459	Accuracy:95.281%
Epoch : 3	[6400/60000 (11%)]	Loss: 0.055648	Accuracy:95.134%
Epoch : 3	[8000/60000 (13%)]	Loss: 0.032606	Accuracy:95.219%
Epoch : 3	[9600/60000 (16%)]	Loss: 0.057986	Accuracy:94.975%
Epoch : 3	[11200/60000 (19%)]	Loss: 0.275520	Accuracy:95.103%
Epoch : 3	[12800/60000 (21%)]	Loss: 0.070716	Accuracy:94.950%
Epoch : 3	[14400/60000 (24%)]	Loss: 0.027030	Accuracy:94.928%
Epoch : 3	[16000/60000 (27%)]	Loss: 0.066866	Accuracy:94.941%
Epoch : 3	[17600/60000 (29%)]	Loss: 0.028506	Accuracy:94.941%
Epoch : 3	[19200/60000 (32%)]	Loss: 0.219603	Accuracy:94.998%
Epoch : 3	[20800/60000 (35%)]	Loss: 0.051435	Accuracy:94.993%
Epoch : 3	[22400/60000 (37%)]	Loss: 0.011864	Accuracy:95.047%
Epoch : 3	[24000/60000 (40%)]	Loss: 0.177396	Accuracy:95.073%
Epoch : 3	[25600/60000 (43%)]	Loss: 0.050125	Accuracy:95.154%
Epoch : 3	[27200/60000 (45%)]	Loss: 0.215653	Accuracy:95.131%
Epoch : 3	[28800/60000 (48%)]	Loss: 0.357761	Accuracy:95.106%
Epoch : 3	[30400/60000 (51%)]	Loss: 0.303260	Accuracy:95.127%
Epoch : 3	[32000/60000 (53%)]	Loss: 0.435719	Accuracy:95.042%
Epoch : 3	[33600/60000 (56%)]	Loss: 0.038899	Accuracy:95.008%
Epoch : 3	[35200/60000 (59%)]	Loss: 0.157955	Accuracy:95.084%
Epoch : 3	[36800/60000 (61%)]	Loss: 0.065332	Accuracy:95.127%
Epoch : 3	[38400/60000 (64%)]	Loss: 0.087521	Accuracy:95.145%
Epoch : 3	[40000/60000 (67%)]	Loss: 0.119639	Accuracy:95.131%
Epoch : 3	[41600/60000 (69%)]	Loss: 0.089203	Accuracy:95.143%
Epoch : 3	[43200/60000 (72%)]	Loss: 0.412446	Accuracy:95.152%
Epoch : 3	[44800/60000 (75%)]	Loss: 0.103714	Accuracy:95.211%
Epoch : 3	[46400/60000 (77%)]	Loss: 0.355738	Accuracy:95.202%
Epoch : 3	[48000/60000 (80%)]	Loss: 0.139295	Accuracy:95.191%
Epoch : 3	[49600/60000 (83%)]	Loss: 0.194295	Accuracy:95.189%
Epoch : 3	[51200/60000 (85%)]	Loss: 0.083338	Accuracy:95.181%
Epoch : 3	[52800/60000 (88%)]	Loss: 0.779976	Accuracy:95.192%
Epoch : 3	[54400/60000 (91%)]	Loss: 0.033422	Accuracy:95.196%
Epoch : 3	[56000/60000 (93%)]	Loss: 0.117889	Accuracy:95.208%
Epoch : 3	[57600/60000 (96%)]	Loss: 0.146757	Accuracy:95.232%
Epoch : 3	[59200/60000 (99%)]	Loss: 0.007108	Accuracy:95.286%
Epoch : 4	[0/60000 (0%)]	Loss: 0.037586	Accuracy:96.875%
Epoch : 4	[1600/60000 (3%)]	Loss: 0.532983	Accuracy:95.282%
Epoch : 4	[3200/60000 (5%)]	Loss: 0.141879	Accuracy:95.575%
Epoch : 4	[4800/60000 (8%)]	Loss: 0.068667	Accuracy:95.882%
Epoch : 4	[6400/60000 (11%)]	Loss: 0.088112	Accuracy:95.740%
Epoch : 4	[8000/60000 (13%)]	Loss: 0.112370	Accuracy:95.754%
Epoch : 4	[9600/60000 (16%)]	Loss: 0.098440	Accuracy:95.650%

Epoch : 4	[11200/60000 (19%)]	Loss: 0.284854	Accuracy:95.593%
Epoch : 4	[12800/60000 (21%)]	Loss: 0.082128	Accuracy:95.605%
Epoch : 4	[14400/60000 (24%)]	Loss: 0.066085	Accuracy:95.468%
Epoch : 4	[16000/60000 (27%)]	Loss: 0.137398	Accuracy:95.422%
Epoch : 4	[17600/60000 (29%)]	Loss: 0.083179	Accuracy:95.389%
Epoch : 4	[19200/60000 (32%)]	Loss: 0.296649	Accuracy:95.372%
Epoch : 4	[20800/60000 (35%)]	Loss: 0.079538	Accuracy:95.387%
Epoch : 4	[22400/60000 (37%)]	Loss: 0.021219	Accuracy:95.399%
Epoch : 4	[24000/60000 (40%)]	Loss: 0.045635	Accuracy:95.460%
Epoch : 4	[25600/60000 (43%)]	Loss: 0.025185	Accuracy:95.431%
Epoch : 4	[27200/60000 (45%)]	Loss: 0.313549	Accuracy:95.384%
Epoch : 4	[28800/60000 (48%)]	Loss: 0.358900	Accuracy:95.384%
Epoch : 4	[30400/60000 (51%)]	Loss: 0.290802	Accuracy:95.354%
Epoch : 4	[32000/60000 (53%)]	Loss: 0.031017	Accuracy:95.367%
Epoch : 4	[33600/60000 (56%)]	Loss: 0.085758	Accuracy:95.362%
Epoch : 4	[35200/60000 (59%)]	Loss: 0.152539	Accuracy:95.416%
Epoch : 4	[36800/60000 (61%)]	Loss: 0.162986	Accuracy:95.441%
Epoch : 4	[38400/60000 (64%)]	Loss: 0.174990	Accuracy:95.444%
Epoch : 4	[40000/60000 (67%)]	Loss: 0.521610	Accuracy:95.456%
Epoch : 4	[41600/60000 (69%)]	Loss: 0.122210	Accuracy:95.484%
Epoch : 4	[43200/60000 (72%)]	Loss: 0.318961	Accuracy:95.503%
Epoch : 4	[44800/60000 (75%)]	Loss: 0.035795	Accuracy:95.550%
Epoch : 4	[46400/60000 (77%)]	Loss: 0.384703	Accuracy:95.535%
Epoch : 4	[48000/60000 (80%)]	Loss: 0.491851	Accuracy:95.561%
Epoch : 4	[49600/60000 (83%)]	Loss: 0.055732	Accuracy:95.561%
Epoch : 4	[51200/60000 (85%)]	Loss: 0.087599	Accuracy:95.583%
Epoch : 4	[52800/60000 (88%)]	Loss: 0.285666	Accuracy:95.597%
Epoch : 4	[54400/60000 (91%)]	Loss: 0.043655	Accuracy:95.602%
Epoch : 4	[56000/60000 (93%)]	Loss: 0.363469	Accuracy:95.604%
Epoch : 4	[57600/60000 (96%)]	Loss: 0.103343	Accuracy:95.624%
Epoch : 4	[59200/60000 (99%)]	Loss: 0.018067	Accuracy:95.676%

0.3 SVM Classifier for MNIST DATA SET

```
[0]: from keras.datasets import mnist
      (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
[0]: x_train = x_train.reshape(60000,784)
      x_test = x_test.reshape(10000,784)
```

```
[0]: ## Applying HOG feature extraction
```

```
[0]: from sklearn.svm import SVC
      from sklearn.metrics import accuracy_score
      from skimage.feature import hog
      from sklearn import preprocessing
```

```
from collections import Counter
```

```
[38]: list_hog_train = []  
for feature in x_train:  
    fd = hog(feature.reshape((28,28)), orientations=10,  
    ↪pixels_per_cell=(7,7),cells_per_block=(1,1),visualize=False )  
    list_hog_train.append(fd)  
hog_features = np.array(list_hog_train, 'float64')  
preProcess = preprocessing.MaxAbsScaler().fit(hog_features)  
hog_features_transformed_train = preProcess.transform(hog_features)  
print(hog_features_transformed_train.shape)
```

(60000, 160)

```
[0]: ## Extracting hog feature for test data
```

```
[41]: list_hog_test = []  
for feature in x_test:  
    fd = hog(feature.reshape((28,28)), orientations=10,  
    ↪pixels_per_cell=(7,7),cells_per_block=(1,1),visualize=False )  
    list_hog_test.append(fd)  
hog_features_test = np.array(list_hog_test, 'float64')  
preProcess = preprocessing.MaxAbsScaler().fit(hog_features_test)  
hog_features_transformed_test = preProcess.transform(hog_features_test)  
print(hog_features_transformed_test.shape)
```

(10000, 160)

```
[0]: model = SVC()  
model.fit(hog_features_transformed_train,y_train)  
y_pred = model.predict(hog_features_transformed_test)
```

Fitting the model with best parameters

```
[43]: print(accuracy_score(y_test, y_pred))
```

0.9713