

q1

April 3, 2020

1 Question 1, PCA ANALYSIS

```
[40]: url="/home/abhishek/dev/Semester_2/SMAI/Assignments/Assignment_3/dataset"
```

```
[41]: import numpy as np
import pandas as pd
import os
import matplotlib.pyplot as plt
from skimage import transform, io
import warnings
warnings.filterwarnings('ignore')
from mpl_toolkits.mplot3d import Axes3D
```

```
[42]: class PCA:

    def read_image(self,url):
        print("Reading Image from directory...")
        dir_list = os.listdir(url)
        faces = []
        image_labels = []
        for image in dir_list:
            labels = image.split("_")
            image_labels.append(labels[0])

        for image in dir_list:
            img = io.imread(url+"/"+image)
            img = img.astype(np.uint8)
            # converting to grayscale
            rgb_weights = [0.2989, 0.5870, 0.1141]
            grayscale_image = np.dot(img[...,:3], rgb_weights)
            #normalizing image
            small_grey = transform.resize(grayscale_image, (64,64),
↪mode='symmetric', preserve_range=True)
            reshape_img = small_grey.reshape(1, 4096)
            faces.append(reshape_img[0])
        X_train = np.asarray(faces)
```

```

        # print(self.X_train[0])
        print(X_train.shape)
        print("Reading Image from directory Done...")
        return image_labels,X_train

def apply_pca(self,X):
    print("Applying PCA.....")
    eig_val, eig_mat = np.linalg.eig(np.cov(X))
    idx = eig_val.argsort()[::-1]
    eig_val = eig_val[idx]
    eig_mat = eig_mat[:,idx]
    print("Eigen Matrix calculation done..")
    # taking principal components
    M = []
    print(eig_mat.shape)
    full_pc = eig_mat.shape[0]
    for numpc in range(0,full_pc,10):
        eigen_coeff = eig_mat[:,range(numpc)]
        score_mat = np.dot(eigen_coeff.T,X)
        final_score = np.dot(eigen_coeff,score_mat).T
        # print(final_score.shape)
        val = np.linalg.norm(X.T-final_score,'fro')
        M.append(val)
        # print(val)

    mse = np.asarray(M)
    print("Applying PCA done.....")
    return mse,eig_mat,eig_val

def plot_mse_vs_principal_component(self,mse,eig_mat):
    full_pc = eig_mat.shape[1]
    plt.figure()
    plt.plot(range(0,full_pc,10),mse,'r')
    plt.show()

def no_of_component_such_that_mse_less_than_20(self,eig_mat,X):
#     N = 20 #from observation no of principal component = 50
    eigen_coeff = eig_mat[:,range(24)]
    score_mat = np.dot(eigen_coeff.T,X)
    final = np.dot(eigen_coeff,score_mat).T
    final = final*255
    final_score = final.astype(int)
    fig,axes = plt.subplots(4,5,figsize=(9,9),subplot_kw={'xticks':[],'y
    ticks':[]},gridspec_kw=dict(hspace=0.01, wspace=0.01))
    for i, ax in enumerate(axes.flat):
        ax.imshow(final_score[i].reshape(64,64),cmap='gray')

```

```

plt.show()
return final_score

def scatter_plot_pca(self,X_new,X):
    #1-D plot
    val=0
    plt.plot(X[:,0],np.zeros_like(X[:,0])+val,'o')
    plt.plot(X_new[:,0],np.zeros_like(X_new[:,0])+val,"x")
    plt.show()

    # 2-D plot
    plt.scatter(X[:,0],X[:,1], alpha=0.2)
    plt.scatter(X_new[:,0],X_new[:,1], alpha=0.8)#=self.image_labels)
    plt.show()

    #3-D plot
    fig = plt.figure()
    ax = fig.add_subplot(111, projection='3d')
    ax.scatter(X[:,0], X[:,1], X[:,2], marker='o')
    ax.scatter(X_new[:,0], X_new[:,1], X_new[:,2], marker='^')
    ax.set_xlabel('X Label')
    ax.set_ylabel('Y Label')
    ax.set_zlabel('Z Label')
    plt.show()

```

```

[43]: pca = PCA()
      image_label,X_train = pca.read_image(url)
      #normalization
      X = X_train
      X_train = X_train/255
      #applying PCA
      mse,eig_mat,eig_val = pca.apply_pca(X_train.T)

```

```


Reading Image from directory...
(520, 4096)
Reading Image from directory Done...
Applying PCA...
Eigen Matrix calculation done..
(4096, 4096)
Applying PCA done...

```

```

[44]: pca.plot_mse_vs_principal_component(mse,eig_mat)

```



```
[45]: final_score = pca.no_of_component_such_that_mse_less_than_20(eig_mat,X_train.T)
```



```
[46]: ## Check N has Mean Squared error has less than 20%
```



```
[47]: eigen_val_n = eig_val[0:24]
      x = np.sum(eigen_val_n)
      y = np.sum(eig_val)
      print(x/y)
```


```
(0.804354886673874-5.750783086893538e-33j)
```

```
[48]: ## for N = 24, mean square error is less than 20%
```

1.1 `pca.scatter_plot_pca(final_score,X)`

```
[49]: pca.scatter_plot_pca(final_score,X)
```





[]:

q2

April 3, 2020

1 Q2 - Logistic Regression After applying PCA

```
[15]: import os
import sys
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from skimage import transform, io
from sklearn.metrics import accuracy_score

url_train_csv = "/home/abhishek/dev/Semester_2/SMAI/Assignments/Assignment_3/q2/
↳sample_train_2.txt"
url_test_csv = "/home/abhishek/dev/Semester_2/SMAI/Assignments/Assignment_3/q2/
↳sample_test_2.txt"
sample_output_csv = "/home/abhishek/dev/Semester_2/SMAI/Assignments/
↳Assignment_3/q2/output_of_sample_test_2.txt"

class LogisticRegression:
    def read_csv_file(self, marker):
        image_files = []
        f = None
        if(marker == "train"):
            f = open(url_train_csv, 'r')
        elif(marker == "test"):
            f = open(url_test_csv, 'r')
        else:
            f = open(sample_output_csv, 'r')

        lines = f.readlines()
        for line in lines:
            image_files.append(line.strip())
        return image_files

    def split_image_dir_and_label(self, images_directory, marker):
        if(marker == "train"):
            # print("splitting directory and labels")
            label = []
```



```

        directory = []
        for l in images_directory:
            x = l.split(" ")
            label.append(x[1])
            directory.append(x[0])
            # print("splitting directory and labels done ....")
        return label, directory
    elif(marker == "test"):
        # print("read lines...")
        lines = []
        for l in images_directory:
            lines.append(l)
            # print("reading lines done...")
        return lines

def read_image_from_directory_to_grayscale(self,directory):
    # print("reading images from directory and downscaling images")
    faces = []
    for i in directory:
        img = io.imread(i)
        img = img.astype(np.uint8)
        # converting to grayscale
        rgb_weights = [0.2989, 0.5870, 0.1141]
        grayscale_image = np.dot(img[...,:3], rgb_weights)
        small_grey = transform.resize(grayscale_image, (64,64),
↪mode='symmetric', preserve_range=True)
        reshape_img = small_grey.reshape(1, 4096)
        faces.append(reshape_img[0])
    X = np.asarray(faces)
    # print("reading images from directory and downscaling images done...")
    return X

def apply_pca(self,X):
    # print("Applying PCA on the image set")
    eig_val, eig_mat = np.linalg.eig(np.cov(X))
    idx = eig_val.argsort()[::-1]
    eig_val = eig_val[idx]
    eig_mat = eig_mat[:,idx]
    eigen_coeff = eig_mat[:,range(50)]
    X_PCA = np.dot(eigen_coeff.T,X)
    # print("Applying PCA on the image set done .....")
    return X_PCA.T

def sigmoid(self,z):
    return 1/(1+np.exp(-z))

def cost(self,w,b,X,y,lmd=10):

```

```

        # print("Calculating Cost")
        m = X.shape[0]
        z = np.matmul(X,w) + b
        hx = self.sigmoid(z)
        J = (-1/m)*( np.sum( y * np.log(hx) + (1. - y) * np.log(1.- hx) ) )
        J += (lmd/(2*m))* np.matmul(w,w)
        # print("Calculating Cost done ..")
        return J

    def gradient_descent(self,w,b,X,y,learning_rate=0.
→01,lmd=10,no_of_iteration=10000):
        # print("Applying Gradient Descent")
        m = X.shape[0]
        # print("Initial cost: {}".format( self.cost(w,b,X,y) ))
        for i in range(no_of_iteration):
            z = np.matmul(X,w) + b
            hx = self.sigmoid(z)
            dw = (1/m)*np.matmul(X.T,hx-y)
            db = (1/m)*np.sum(hx-y)
            factor = 1-( (learning_rate * lmd)/m)
            w = w*factor - learning_rate*dw
            b = b - learning_rate*db
            # if i % 500 == 0:
                # print("Iteration {} cost :{}".format(i,self.cost(w,b,X,y)))
        # print("Final cost: {}".format( self.cost(w,b,X,y) ))
        # print("Applying Gradient Descent done ....")
        return w,b

    def accuracy(self,w,b,X,y):
        # print("Calculating Accuracy")
        m = X.shape[0]
        z = np.matmul(X,w) + b
        hx = self.sigmoid(z)
        pred = np.round(hx)
        correct_pred = (pred==y)
        total = np.sum(correct_pred)
        # print("Calculating Accuracy done ....." )
        return (total*100)/m

    def test(self,w,b,X):
        # m = X.shape[0]
        z = np.matmul(X,w)+b
        hx = self.sigmoid(z)
        pred = np.round(hx)
        return pred

lr = LogisticRegression()

```

```
[ ]: ## Applying PCA on the above dataset
```

```
[2]: lines = lr.read_csv_file("train")
      # print(images_directory)
      label, directory = lr.split_image_dir_and_label(lines, "train")
      # label = np.asarray(label)
      # print(label)
      unique_labels = np.unique(np.asarray(label))
      # print(unique_labels)
      X_train = lr.read_image_from_directory_to_grayscale(directory)
      # normalizing
      X_train = X_train/255
      #applying PCA
      X_PCA = lr.apply_pca(X_train.T)
      # print(X_PCA.shape)
      # print(X_PCA[0])
      m = X_train.shape[0]
```

```
[3]: print(X_PCA.shape)
```

(520, 50)

```
[7]: ## Training Multi-Class Classifier for each unique class in Tranining set
```

```
[4]: weight_label_map = {}
      i = 0
      THETA = []
      BIAS = []

      for unique in unique_labels:
          y_train = []
          for l in label:
              if(unique == l):
                  y_train.append(1)
              else:
                  y_train.append(0)

          weight_label_map[i] = unique
          i = i+1
          y_train = np.asarray(y_train)
          # print(y_train)
          w = np.zeros(X_PCA.shape[1], dtype=np.float64)
          b = 0.0
          w,b = lr.gradient_descent(w,b,X_PCA,y_train)
          THETA.append(w)
          BIAS.append(b)
```

```
[5]: ## ALL Unique Labels in Training set
```

```
[6]: print(weight_label_map)
```

```
{0: '000', 1: '001', 2: '002', 3: '003', 4: '004', 5: '005', 6: '006', 7: '007'}
```

```
[12]: ## Applying PCA on Test Data set
```

```
[7]: THETA = np.asarray(THETA)
BIAS = np.asarray(BIAS)

lr1 = LogisticRegression()
lines = lr1.read_csv_file("test")
directory = lr1.split_image_dir_and_label(lines,"test")
X_test = lr1.read_image_from_directory_to_grayscale(directory)
X_test = X_test/255
X_Test_PCA = lr1.apply_pca(X_test.T)
# print(X_Test_PCA.shape)
```

```
[8]: ## Predicting The labels of Test Data by taking maximum sigmoid value of label_  
↪ along all classifier
```

```
[9]: y_pred = []
for X in X_Test_PCA:
    i = 0
    max_hx = 0.0
    index = 0
    for w,b in zip(THETA,BIAS):
        pred = lr1.test(w,b,X)
        if(pred > max_hx):
            max_hx = pred
            index = i
        i = i+1
    y_pred.append(weight_label_map[index])

y_pred = np.asarray(y_pred)
for prediction in y_pred:
    print(prediction)
```

000

000

000

002

001

002

003

005

000
000
000
002
000
002
000
000
007
000
000
002
007
000
000
000
000
000
000
006
006
001
006
000
000
000
003
006
002
007
000
002
000
000
000
000
004
000
000
003
003
001
000
006
006
001
005
000
000
001

002
000
000
003
002
001
000
000
000
004
000
000
004
007
006
006
007
007
000
000
000
006
000
006
000
006
004
007
000
007
007
002
000
003
005
002
000
003
000
005
000
003
001
007
000
000
005
000

007
000
000
000
000
000
000
000
000
007
002
003
007
000
000
000
003
007
001
000
000
000
007
004
000
000
002
000
002
006
000
005
000
000
001
000
004
000
005
003
003
000
000
006
007
001
000
000

005
000
000
000
004
000
007
005
000
000
000
002
000
000
004
003
001
000
005
000
004
000
000
000
006
000
000
003
000
000
000
000
000
005
005
000
005
007
000
000
000
002
006
007
000
007
004
003

005
002
000
003
000
000
004
000
003
002
004
004
000
000
006
007
000
004
000
001
003
000
000
000
007
003
002
000
003
000
000
000
000
000
000
005
002
000
006
000
000
005
000
004
000
000
002
005
002

000
000
000
000
000
000
000
002
000
000
005
004
005
000
000
001
007
007
001
000
000
003
003
000
000
000
004
003
007
003
000
002
007
000
007
005
000
000
000
006
004
000
000
007
000
000
007
000

003
000
001
005
005
000
000
005
004
002
000
000
000
004
006
004
000
005
004
007
002
000
007
000
005
002
000
007
005
000
007
002
007
004
003
000
000
000
003
002
000
000
000
003
000
001
003
000

000
001
000
000
003
000
000
000
007
000
004
007
000
000
004
000
002
001
000
000
000
000
000
000
000
007
000
000
000
004
000
000
007
000
000
002
000
005
000
000
000
000
000
003
005
000
004
000
002

000
000
000
000
005
004
007
000
002
000
000
004
003
005
007
004
004
001
006
004
004
007
005
007
000
000
003
007
000
004
000
004
003
004
007
002
005
003
006
002
000
000
002
003
005
000
004
000

000
000
003
005
007
007
000
000
000
000
000
003
000
002
002
006
000
003
003
003
001
000
000
000
006
000
003
007
001
000
000
000
004
000
002
000
005
000
000
000
006
002
004
000
000
001
007
005

007
000
004
003
006
000
000
000
004
005
001
002
002
007
003
000
004
007
000
000
000
003
004
004
002
000
000
000
002
001
007
003

```
[16]: labels = lr.read_csv_file("label")
```

```
[18]: print(len(labels))
```

520

```
[20]: labels = np.asarray(labels)
```

```
[21]: print(labels.shape)
```

(520,)

```
[ ]: ## Accuracy Score for predicting the labels
```

```
[22]: accuracy_score(labels,y_pred)
```

[22] : 0.5846153846153846

[] :

q3

April 3, 2020

0.1 Question 3 - MNIST Classification using PyTorch

```
[0]: import torch
import numpy as np
```

Load Data Sets

```
[2]: from torchvision import datasets
import torchvision.transforms as transforms

# number of subprocesses to use for data loading
num_workers = 0
# how many samples per batch to load
batch_size = 20

# convert data to torch.FloatTensor
transform = transforms.ToTensor()

# choose the training and test datasets
train_data = datasets.MNIST(root='data', train=True,
                             download=True, transform=transform)
test_data = datasets.MNIST(root='data', train=False,
                             download=True, transform=transform)

# prepare data loaders
train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                             num_workers=num_workers)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,
                                             num_workers=num_workers)
```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to
data/MNIST/raw/train-images-idx3-ubyte.gz

HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))

Extracting data/MNIST/raw/train-images-idx3-ubyte.gz to data/MNIST/raw
Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to
data/MNIST/raw/train-labels-idx1-ubyte.gz

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

Extracting data/MNIST/raw/train-labels-idx1-ubyte.gz to data/MNIST/raw
Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to
data/MNIST/raw/t10k-images-idx3-ubyte.gz

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

Extracting data/MNIST/raw/t10k-images-idx3-ubyte.gz to data/MNIST/raw
Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to
data/MNIST/raw/t10k-labels-idx1-ubyte.gz

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```


Extracting data/MNIST/raw/t10k-labels-idx1-ubyte.gz to data/MNIST/raw
Processing...
Done!

Visualize a batch of Training Set

```
[3]: import matplotlib.pyplot as plt
      %matplotlib inline

      # obtain one batch of training images
      dataiter = iter(train_loader)
      images, labels = dataiter.next()
      images = images.numpy()

      # plot the images in the batch, along with the corresponding labels
      fig = plt.figure(figsize=(25, 4))
      for idx in np.arange(20):
          ax = fig.add_subplot(2, 20/2, idx+1, xticks=[], yticks=[])
          ax.imshow(np.squeeze(images[idx]), cmap='gray')
          # print out the correct label for each image
          # .item() gets the value contained in a Tensor
          ax.set_title(str(labels[idx].item()))
```



Define the Network Architecture

```
[0]: import torch.nn as nn
import torch.nn.functional as F

## Define the NN architecture
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(28 * 28, 512)
        # linear layer (n_hidden -> hidden_2)
        self.fc2 = nn.Linear(512, 512)
        # linear layer (n_hidden -> 10)
        self.fc3 = nn.Linear(512, 10)
        # dropout layer (p=0.2)
        # dropout prevents overfitting of data
        self.dropout = nn.Dropout(0.2)

    def forward(self, x):
        # flatten image input
        x = x.view(-1, 28 * 28)
        # add hidden layer, with relu activation function
        x = F.relu(self.fc1(x))
        return x
```

```
[5]: # initialize the NN
model = Net()
print(model)
```

```
Net(
  (fc1): Linear(in_features=784, out_features=512, bias=True)
  (fc2): Linear(in_features=512, out_features=512, bias=True)
  (fc3): Linear(in_features=512, out_features=10, bias=True)
  (dropout): Dropout(p=0.2, inplace=False)
)
```

Specify Loss Function and Optimizer

```
[0]: criterion = nn.CrossEntropyLoss()

# specify optimizer
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
```

Train the Network

```

[7]: n_epochs = 30 # suggest training between 20-50 epochs

model.train() # prep model for training

for epoch in range(n_epochs):
    # monitor training loss
    train_loss = 0.0

    #####
    # train the model #
    #####
    for data, target in train_loader:
        # clear the gradients of all optimized variables
        optimizer.zero_grad()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # backward pass: compute gradient of the loss with respect to model
        ↪ parameters
        loss.backward()
        # perform a single optimization step (parameter update)
        optimizer.step()
        # update running training loss
        train_loss += loss.item()*data.size(0)

    # print training statistics
    # calculate average loss over an epoch
    train_loss = train_loss/len(train_loader.dataset)

    print('Epoch: {} \tTraining Loss: {:.6f}'.format(
        epoch+1,
        train_loss
    ))

```

```

Epoch: 1      Training Loss: 0.804222
Epoch: 2      Training Loss: 0.411707
Epoch: 3      Training Loss: 0.370065
Epoch: 4      Training Loss: 0.348959
Epoch: 5      Training Loss: 0.335502
Epoch: 6      Training Loss: 0.325918
Epoch: 7      Training Loss: 0.318626
Epoch: 8      Training Loss: 0.312826
Epoch: 9      Training Loss: 0.308064
Epoch: 10     Training Loss: 0.304058
Epoch: 11     Training Loss: 0.300624
Epoch: 12     Training Loss: 0.297635

```

Epoch: 13	Training Loss: 0.294999
Epoch: 14	Training Loss: 0.292651
Epoch: 15	Training Loss: 0.290539
Epoch: 16	Training Loss: 0.288626
Epoch: 17	Training Loss: 0.286881
Epoch: 18	Training Loss: 0.285279
Epoch: 19	Training Loss: 0.283802
Epoch: 20	Training Loss: 0.282434
Epoch: 21	Training Loss: 0.281160
Epoch: 22	Training Loss: 0.279970
Epoch: 23	Training Loss: 0.278856
Epoch: 24	Training Loss: 0.277808
Epoch: 25	Training Loss: 0.276820
Epoch: 26	Training Loss: 0.275887
Epoch: 27	Training Loss: 0.275002
Epoch: 28	Training Loss: 0.274163
Epoch: 29	Training Loss: 0.273365
Epoch: 30	Training Loss: 0.272604

Test the Trained Network

```
[8]: test_loss = 0.0
class_correct = list(0. for i in range(10))
class_total = list(0. for i in range(10))

model.eval() # prep model for *evaluation*

for data, target in test_loader:
    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data)
    # calculate the loss
    loss = criterion(output, target)
    # update test loss
    test_loss += loss.item()*data.size(0)
    # convert output probabilities to predicted class
    _, pred = torch.max(output, 1)
    # compare predictions to true label
    correct = np.squeeze(pred.eq(target.data.view_as(pred)))
    # calculate test accuracy for each object class
    for i in range(batch_size):
        label = target.data[i]
        class_correct[label] += correct[i].item()
        class_total[label] += 1

# calculate and print avg test loss
test_loss = test_loss/len(test_loader.dataset)
print('Test Loss: {:.6f}\n'.format(test_loss))
```

```

for i in range(10):
    if class_total[i] > 0:
        print('Test Accuracy of %5s: %2d%% (%2d/%2d)' % (
            str(i), 100 * class_correct[i] / class_total[i],
            np.sum(class_correct[i]), np.sum(class_total[i])))
    else:
        print('Test Accuracy of %5s: N/A (no training examples)' % (classes[i]))

print('\nTest Accuracy (Overall): %2d%% (%2d/%2d)' % (
    100. * np.sum(class_correct) / np.sum(class_total),
    np.sum(class_correct), np.sum(class_total)))

```

Test Loss: 0.274007

```

Test Accuracy of    0: 98% (962/980)
Test Accuracy of    1: 97% (1109/1135)
Test Accuracy of    2: 88% (918/1032)
Test Accuracy of    3: 90% (914/1010)
Test Accuracy of    4: 92% (909/982)
Test Accuracy of    5: 87% (781/892)
Test Accuracy of    6: 94% (910/958)
Test Accuracy of    7: 92% (946/1028)
Test Accuracy of    8: 88% (864/974)
Test Accuracy of    9: 90% (916/1009)

```

Test Accuracy (Overall): 92% (9229/10000)

Visualize Sample Test Results


```

[9]: dataiter = iter(test_loader)
     images, labels = dataiter.next()

     # get sample outputs
     output = model(images)
     # convert output probabilities to predicted class
     _, preds = torch.max(output, 1)
     # prep images for display
     images = images.numpy()

     # plot the images in the batch, along with predicted and true labels
     fig = plt.figure(figsize=(25, 4))
     for idx in np.arange(20):
         ax = fig.add_subplot(2, 20/2, idx+1, xticks=[], yticks=[])
         ax.imshow(np.squeeze(images[idx]), cmap='gray')
         ax.set_title("{} ({}).format(str(preds[idx].item()), str(labels[idx].
         →item()))),
         color=("green" if preds[idx]==labels[idx] else "red"))

```



0.2 MNIST Classification using CNN

```
[0]: import numpy as np # to handle matrix and data operation
import pandas as pd # to read csv and handle dataframe

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.utils.data
from torch.autograd import Variable

from sklearn.model_selection import train_test_split
```

```
[0]: from keras.datasets import mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
[0]: x_train = x_train.reshape(60000,784)
x_test = x_test.reshape(10000,784)
```

```
[0]: BATCH_SIZE = 32
```

```
[0]: torch_X_train = torch.from_numpy(x_train).type(torch.LongTensor)
torch_y_train = torch.from_numpy(y_train).type(torch.LongTensor)
```

```
[0]: torch_X_test = torch.from_numpy(x_test).type(torch.LongTensor)
torch_y_test = torch.from_numpy(y_test).type(torch.LongTensor)
```

```
[0]: train = torch.utils.data.TensorDataset(torch_X_train,torch_y_train)
test = torch.utils.data.TensorDataset(torch_X_test,torch_y_test)
```

```
[0]: train_loader = torch.utils.data.DataLoader(train, batch_size = BATCH_SIZE,
    ↪shuffle = False)
test_loader = torch.utils.data.DataLoader(test, batch_size = BATCH_SIZE,
    ↪shuffle = False)
```

```
[59]: class MLP(nn.Module):
    def __init__(self):
        super(MLP, self).__init__()
        self.linear1 = nn.Linear(784,250)
        self.linear2 = nn.Linear(250,100)
        self.linear3 = nn.Linear(100,10)

    def forward(self,X):
        X = F.relu(self.linear1(X))
        X = F.relu(self.linear2(X))
        X = self.linear3(X)
        return F.log_softmax(X, dim=1)

mlp = MLP()
print(mlp)
```

```
MLP(
  (linear1): Linear(in_features=784, out_features=250, bias=True)
  (linear2): Linear(in_features=250, out_features=100, bias=True)
  (linear3): Linear(in_features=100, out_features=10, bias=True)
)
```

```
[0]: def fit(model, train_loader):
    optimizer = torch.optim.Adam(model.parameters())#,lr=0.001, betas=(0.9,0.
    ↪999))
    error = nn.CrossEntropyLoss()
    EPOCHS = 5
    model.train()
    for epoch in range(EPOCHS):
        correct = 0
        for batch_idx, (X_batch, y_batch) in enumerate(train_loader):
            var_X_batch = Variable(X_batch).float()
            var_y_batch = Variable(y_batch)
            optimizer.zero_grad()
            output = model(var_X_batch)
            loss = error(output, var_y_batch)
            loss.backward()
            optimizer.step()

            # Total correct predictions
            predicted = torch.max(output.data, 1)[1]
            correct += (predicted == var_y_batch).sum()
            #print(correct)
            if batch_idx % 50 == 0:
                print('Epoch : {} [{} / {}] {:.0f}%]\tLoss: {:.6f}\t Accuracy:{:.
                ↪3f}%'.format(
```



```

epoch, batch_idx*len(X_batch), len(train_loader.dataset),
↪100.*batch_idx / len(train_loader), (loss.data), float(correct*100) /
↪float(BATCH_SIZE*(batch_idx+1)))

```

```
[67]: fit(mlp, train_loader)
```

Epoch : 0	[0/60000 (0%)]	Loss: 0.028779	Accuracy:100.000%
Epoch : 0	[1600/60000 (3%)]	Loss: 0.135546	Accuracy:96.875%
Epoch : 0	[3200/60000 (5%)]	Loss: 0.161492	Accuracy:97.061%
Epoch : 0	[4800/60000 (8%)]	Loss: 0.029837	Accuracy:97.144%
Epoch : 0	[6400/60000 (11%)]	Loss: 0.251311	Accuracy:97.233%
Epoch : 0	[8000/60000 (13%)]	Loss: 0.288707	Accuracy:96.987%
Epoch : 0	[9600/60000 (16%)]	Loss: 0.001341	Accuracy:96.564%
Epoch : 0	[11200/60000 (19%)]	Loss: 0.147621	Accuracy:96.270%
Epoch : 0	[12800/60000 (21%)]	Loss: 0.162637	Accuracy:96.213%
Epoch : 0	[14400/60000 (24%)]	Loss: 0.040754	Accuracy:96.182%
Epoch : 0	[16000/60000 (27%)]	Loss: 0.053642	Accuracy:96.214%
Epoch : 0	[17600/60000 (29%)]	Loss: 0.021809	Accuracy:96.257%
Epoch : 0	[19200/60000 (32%)]	Loss: 0.007951	Accuracy:96.339%
Epoch : 0	[20800/60000 (35%)]	Loss: 0.117186	Accuracy:96.376%
Epoch : 0	[22400/60000 (37%)]	Loss: 0.001032	Accuracy:96.438%
Epoch : 0	[24000/60000 (40%)]	Loss: 0.240103	Accuracy:96.451%
Epoch : 0	[25600/60000 (43%)]	Loss: 0.097887	Accuracy:96.497%
Epoch : 0	[27200/60000 (45%)]	Loss: 0.023727	Accuracy:96.500%
Epoch : 0	[28800/60000 (48%)]	Loss: 0.015119	Accuracy:96.528%
Epoch : 0	[30400/60000 (51%)]	Loss: 0.070183	Accuracy:96.537%
Epoch : 0	[32000/60000 (53%)]	Loss: 0.012500	Accuracy:96.544%
Epoch : 0	[33600/60000 (56%)]	Loss: 0.110607	Accuracy:96.530%
Epoch : 0	[35200/60000 (59%)]	Loss: 0.357193	Accuracy:96.540%
Epoch : 0	[36800/60000 (61%)]	Loss: 0.010914	Accuracy:96.549%
Epoch : 0	[38400/60000 (64%)]	Loss: 0.176304	Accuracy:96.555%
Epoch : 0	[40000/60000 (67%)]	Loss: 0.098055	Accuracy:96.530%
Epoch : 0	[41600/60000 (69%)]	Loss: 0.044533	Accuracy:96.539%
Epoch : 0	[43200/60000 (72%)]	Loss: 0.099969	Accuracy:96.551%
Epoch : 0	[44800/60000 (75%)]	Loss: 0.024746	Accuracy:96.576%
Epoch : 0	[46400/60000 (77%)]	Loss: 0.333618	Accuracy:96.550%
Epoch : 0	[48000/60000 (80%)]	Loss: 0.118122	Accuracy:96.565%
Epoch : 0	[49600/60000 (83%)]	Loss: 0.100130	Accuracy:96.555%
Epoch : 0	[51200/60000 (85%)]	Loss: 0.101556	Accuracy:96.549%
Epoch : 0	[52800/60000 (88%)]	Loss: 0.163493	Accuracy:96.555%
Epoch : 0	[54400/60000 (91%)]	Loss: 0.001940	Accuracy:96.550%
Epoch : 0	[56000/60000 (93%)]	Loss: 0.151699	Accuracy:96.552%
Epoch : 0	[57600/60000 (96%)]	Loss: 0.085508	Accuracy:96.577%
Epoch : 0	[59200/60000 (99%)]	Loss: 0.006413	Accuracy:96.607%
Epoch : 1	[0/60000 (0%)]	Loss: 0.066838	Accuracy:96.875%
Epoch : 1	[1600/60000 (3%)]	Loss: 0.038393	Accuracy:95.895%
Epoch : 1	[3200/60000 (5%)]	Loss: 0.088629	Accuracy:96.566%

Epoch : 1	[4800/60000 (8%)]	Loss: 0.073845	Accuracy:96.978%
Epoch : 1	[6400/60000 (11%)]	Loss: 0.323887	Accuracy:97.015%
Epoch : 1	[8000/60000 (13%)]	Loss: 0.080831	Accuracy:97.049%
Epoch : 1	[9600/60000 (16%)]	Loss: 0.028484	Accuracy:97.103%
Epoch : 1	[11200/60000 (19%)]	Loss: 0.217653	Accuracy:97.044%
Epoch : 1	[12800/60000 (21%)]	Loss: 0.015726	Accuracy:97.015%
Epoch : 1	[14400/60000 (24%)]	Loss: 0.031854	Accuracy:96.993%
Epoch : 1	[16000/60000 (27%)]	Loss: 0.102694	Accuracy:97.050%
Epoch : 1	[17600/60000 (29%)]	Loss: 0.314560	Accuracy:97.164%
Epoch : 1	[19200/60000 (32%)]	Loss: 0.043562	Accuracy:97.192%
Epoch : 1	[20800/60000 (35%)]	Loss: 0.078116	Accuracy:97.211%
Epoch : 1	[22400/60000 (37%)]	Loss: 0.010395	Accuracy:97.214%
Epoch : 1	[24000/60000 (40%)]	Loss: 0.008545	Accuracy:97.241%
Epoch : 1	[25600/60000 (43%)]	Loss: 0.087680	Accuracy:97.285%
Epoch : 1	[27200/60000 (45%)]	Loss: 0.016307	Accuracy:97.264%
Epoch : 1	[28800/60000 (48%)]	Loss: 0.070700	Accuracy:97.295%
Epoch : 1	[30400/60000 (51%)]	Loss: 0.283281	Accuracy:97.315%
Epoch : 1	[32000/60000 (53%)]	Loss: 0.124514	Accuracy:97.331%
Epoch : 1	[33600/60000 (56%)]	Loss: 0.077984	Accuracy:97.265%
Epoch : 1	[35200/60000 (59%)]	Loss: 0.018130	Accuracy:97.272%
Epoch : 1	[36800/60000 (61%)]	Loss: 0.063739	Accuracy:97.277%
Epoch : 1	[38400/60000 (64%)]	Loss: 0.262766	Accuracy:97.250%
Epoch : 1	[40000/60000 (67%)]	Loss: 0.040512	Accuracy:97.252%
Epoch : 1	[41600/60000 (69%)]	Loss: 0.088825	Accuracy:97.235%
Epoch : 1	[43200/60000 (72%)]	Loss: 0.161225	Accuracy:97.213%
Epoch : 1	[44800/60000 (75%)]	Loss: 0.050391	Accuracy:97.232%
Epoch : 1	[46400/60000 (77%)]	Loss: 0.577136	Accuracy:97.194%
Epoch : 1	[48000/60000 (80%)]	Loss: 0.154128	Accuracy:97.183%
Epoch : 1	[49600/60000 (83%)]	Loss: 0.199059	Accuracy:97.155%
Epoch : 1	[51200/60000 (85%)]	Loss: 0.163394	Accuracy:97.146%
Epoch : 1	[52800/60000 (88%)]	Loss: 0.163188	Accuracy:97.161%
Epoch : 1	[54400/60000 (91%)]	Loss: 0.007887	Accuracy:97.163%
Epoch : 1	[56000/60000 (93%)]	Loss: 0.123494	Accuracy:97.150%
Epoch : 1	[57600/60000 (96%)]	Loss: 0.084630	Accuracy:97.167%
Epoch : 1	[59200/60000 (99%)]	Loss: 0.041318	Accuracy:97.197%
Epoch : 2	[0/60000 (0%)]	Loss: 0.101924	Accuracy:93.750%
Epoch : 2	[1600/60000 (3%)]	Loss: 0.192111	Accuracy:96.078%
Epoch : 2	[3200/60000 (5%)]	Loss: 0.070167	Accuracy:97.123%
Epoch : 2	[4800/60000 (8%)]	Loss: 0.023145	Accuracy:97.413%
Epoch : 2	[6400/60000 (11%)]	Loss: 0.164375	Accuracy:97.217%
Epoch : 2	[8000/60000 (13%)]	Loss: 0.206731	Accuracy:97.261%
Epoch : 2	[9600/60000 (16%)]	Loss: 0.013814	Accuracy:97.228%
Epoch : 2	[11200/60000 (19%)]	Loss: 0.189980	Accuracy:97.302%
Epoch : 2	[12800/60000 (21%)]	Loss: 0.094958	Accuracy:97.319%
Epoch : 2	[14400/60000 (24%)]	Loss: 0.015213	Accuracy:97.367%
Epoch : 2	[16000/60000 (27%)]	Loss: 0.093988	Accuracy:97.287%
Epoch : 2	[17600/60000 (29%)]	Loss: 0.621408	Accuracy:97.221%
Epoch : 2	[19200/60000 (32%)]	Loss: 0.108640	Accuracy:97.223%

Epoch : 2	[20800/60000 (35%)]	Loss: 0.021036	Accuracy:97.216%
Epoch : 2	[22400/60000 (37%)]	Loss: 0.001247	Accuracy:97.209%
Epoch : 2	[24000/60000 (40%)]	Loss: 0.015268	Accuracy:97.233%
Epoch : 2	[25600/60000 (43%)]	Loss: 0.101386	Accuracy:97.261%
Epoch : 2	[27200/60000 (45%)]	Loss: 0.131324	Accuracy:97.239%
Epoch : 2	[28800/60000 (48%)]	Loss: 0.069168	Accuracy:97.243%
Epoch : 2	[30400/60000 (51%)]	Loss: 0.025075	Accuracy:97.282%
Epoch : 2	[32000/60000 (53%)]	Loss: 0.016453	Accuracy:97.300%
Epoch : 2	[33600/60000 (56%)]	Loss: 0.020328	Accuracy:97.273%
Epoch : 2	[35200/60000 (59%)]	Loss: 0.872610	Accuracy:97.298%
Epoch : 2	[36800/60000 (61%)]	Loss: 0.101088	Accuracy:97.296%
Epoch : 2	[38400/60000 (64%)]	Loss: 0.088246	Accuracy:97.294%
Epoch : 2	[40000/60000 (67%)]	Loss: 0.119826	Accuracy:97.302%
Epoch : 2	[41600/60000 (69%)]	Loss: 0.052212	Accuracy:97.324%
Epoch : 2	[43200/60000 (72%)]	Loss: 0.325394	Accuracy:97.305%
Epoch : 2	[44800/60000 (75%)]	Loss: 0.062563	Accuracy:97.310%
Epoch : 2	[46400/60000 (77%)]	Loss: 0.133759	Accuracy:97.273%
Epoch : 2	[48000/60000 (80%)]	Loss: 0.219807	Accuracy:97.271%
Epoch : 2	[49600/60000 (83%)]	Loss: 0.019319	Accuracy:97.288%
Epoch : 2	[51200/60000 (85%)]	Loss: 0.045573	Accuracy:97.273%
Epoch : 2	[52800/60000 (88%)]	Loss: 0.129453	Accuracy:97.265%
Epoch : 2	[54400/60000 (91%)]	Loss: 0.037412	Accuracy:97.270%
Epoch : 2	[56000/60000 (93%)]	Loss: 0.061451	Accuracy:97.275%
Epoch : 2	[57600/60000 (96%)]	Loss: 0.268131	Accuracy:97.278%
Epoch : 2	[59200/60000 (99%)]	Loss: 0.029599	Accuracy:97.316%
Epoch : 3	[0/60000 (0%)]	Loss: 0.151179	Accuracy:93.750%
Epoch : 3	[1600/60000 (3%)]	Loss: 0.128003	Accuracy:96.446%
Epoch : 3	[3200/60000 (5%)]	Loss: 0.006657	Accuracy:97.308%
Epoch : 3	[4800/60000 (8%)]	Loss: 0.015271	Accuracy:97.392%
Epoch : 3	[6400/60000 (11%)]	Loss: 0.127014	Accuracy:97.404%
Epoch : 3	[8000/60000 (13%)]	Loss: 0.137492	Accuracy:97.323%
Epoch : 3	[9600/60000 (16%)]	Loss: 0.011746	Accuracy:97.290%
Epoch : 3	[11200/60000 (19%)]	Loss: 0.193034	Accuracy:97.391%
Epoch : 3	[12800/60000 (21%)]	Loss: 0.151988	Accuracy:97.506%
Epoch : 3	[14400/60000 (24%)]	Loss: 0.001826	Accuracy:97.561%
Epoch : 3	[16000/60000 (27%)]	Loss: 0.038850	Accuracy:97.493%
Epoch : 3	[17600/60000 (29%)]	Loss: 0.005036	Accuracy:97.533%
Epoch : 3	[19200/60000 (32%)]	Loss: 0.041748	Accuracy:97.577%
Epoch : 3	[20800/60000 (35%)]	Loss: 0.061989	Accuracy:97.552%
Epoch : 3	[22400/60000 (37%)]	Loss: 0.000083	Accuracy:97.562%
Epoch : 3	[24000/60000 (40%)]	Loss: 0.010684	Accuracy:97.562%
Epoch : 3	[25600/60000 (43%)]	Loss: 0.035633	Accuracy:97.554%
Epoch : 3	[27200/60000 (45%)]	Loss: 0.003687	Accuracy:97.536%
Epoch : 3	[28800/60000 (48%)]	Loss: 0.059841	Accuracy:97.562%
Epoch : 3	[30400/60000 (51%)]	Loss: 0.055981	Accuracy:97.581%
Epoch : 3	[32000/60000 (53%)]	Loss: 0.057697	Accuracy:97.596%
Epoch : 3	[33600/60000 (56%)]	Loss: 0.099345	Accuracy:97.589%
Epoch : 3	[35200/60000 (59%)]	Loss: 0.004503	Accuracy:97.593%

Epoch : 3	[36800/60000 (61%)]	Loss: 0.132729	Accuracy:97.597%
Epoch : 3	[38400/60000 (64%)]	Loss: 0.078958	Accuracy:97.580%
Epoch : 3	[40000/60000 (67%)]	Loss: 0.049452	Accuracy:97.577%
Epoch : 3	[41600/60000 (69%)]	Loss: 0.001719	Accuracy:97.574%
Epoch : 3	[43200/60000 (72%)]	Loss: 0.175857	Accuracy:97.580%
Epoch : 3	[44800/60000 (75%)]	Loss: 0.034069	Accuracy:97.589%
Epoch : 3	[46400/60000 (77%)]	Loss: 0.021272	Accuracy:97.599%
Epoch : 3	[48000/60000 (80%)]	Loss: 0.064088	Accuracy:97.620%
Epoch : 3	[49600/60000 (83%)]	Loss: 0.003137	Accuracy:97.625%
Epoch : 3	[51200/60000 (85%)]	Loss: 0.014169	Accuracy:97.628%
Epoch : 3	[52800/60000 (88%)]	Loss: 0.221221	Accuracy:97.625%
Epoch : 3	[54400/60000 (91%)]	Loss: 0.001268	Accuracy:97.612%
Epoch : 3	[56000/60000 (93%)]	Loss: 0.063965	Accuracy:97.596%
Epoch : 3	[57600/60000 (96%)]	Loss: 0.254212	Accuracy:97.604%
Epoch : 3	[59200/60000 (99%)]	Loss: 0.000153	Accuracy:97.633%
Epoch : 4	[0/60000 (0%)]	Loss: 0.077204	Accuracy:96.875%
Epoch : 4	[1600/60000 (3%)]	Loss: 0.160978	Accuracy:97.059%
Epoch : 4	[3200/60000 (5%)]	Loss: 0.022128	Accuracy:97.463%
Epoch : 4	[4800/60000 (8%)]	Loss: 0.087849	Accuracy:97.641%
Epoch : 4	[6400/60000 (11%)]	Loss: 0.282174	Accuracy:97.512%
Epoch : 4	[8000/60000 (13%)]	Loss: 0.022924	Accuracy:97.572%
Epoch : 4	[9600/60000 (16%)]	Loss: 0.057789	Accuracy:97.456%
Epoch : 4	[11200/60000 (19%)]	Loss: 0.173873	Accuracy:97.480%
Epoch : 4	[12800/60000 (21%)]	Loss: 0.177891	Accuracy:97.498%
Epoch : 4	[14400/60000 (24%)]	Loss: 0.058999	Accuracy:97.568%
Epoch : 4	[16000/60000 (27%)]	Loss: 0.087360	Accuracy:97.493%
Epoch : 4	[17600/60000 (29%)]	Loss: 0.002993	Accuracy:97.544%
Epoch : 4	[19200/60000 (32%)]	Loss: 0.007460	Accuracy:97.577%
Epoch : 4	[20800/60000 (35%)]	Loss: 0.172938	Accuracy:97.576%
Epoch : 4	[22400/60000 (37%)]	Loss: 0.003647	Accuracy:97.597%
Epoch : 4	[24000/60000 (40%)]	Loss: 0.003558	Accuracy:97.616%
Epoch : 4	[25600/60000 (43%)]	Loss: 0.123521	Accuracy:97.651%
Epoch : 4	[27200/60000 (45%)]	Loss: 0.062108	Accuracy:97.650%
Epoch : 4	[28800/60000 (48%)]	Loss: 0.079932	Accuracy:97.683%
Epoch : 4	[30400/60000 (51%)]	Loss: 0.002716	Accuracy:97.703%
Epoch : 4	[32000/60000 (53%)]	Loss: 0.218044	Accuracy:97.687%
Epoch : 4	[33600/60000 (56%)]	Loss: 0.009926	Accuracy:97.693%
Epoch : 4	[35200/60000 (59%)]	Loss: 0.001960	Accuracy:97.707%
Epoch : 4	[36800/60000 (61%)]	Loss: 0.003491	Accuracy:97.706%
Epoch : 4	[38400/60000 (64%)]	Loss: 0.193213	Accuracy:97.684%
Epoch : 4	[40000/60000 (67%)]	Loss: 0.121122	Accuracy:97.677%
Epoch : 4	[41600/60000 (69%)]	Loss: 0.058635	Accuracy:97.672%
Epoch : 4	[43200/60000 (72%)]	Loss: 0.121910	Accuracy:97.668%
Epoch : 4	[44800/60000 (75%)]	Loss: 0.020887	Accuracy:97.711%
Epoch : 4	[46400/60000 (77%)]	Loss: 0.099990	Accuracy:97.676%
Epoch : 4	[48000/60000 (80%)]	Loss: 0.167490	Accuracy:97.677%
Epoch : 4	[49600/60000 (83%)]	Loss: 0.096079	Accuracy:97.671%
Epoch : 4	[51200/60000 (85%)]	Loss: 0.009618	Accuracy:97.687%

```
Epoch : 4 [52800/60000 (88%)]    Loss: 0.322359    Accuracy:97.714%
Epoch : 4 [54400/60000 (91%)]    Loss: 0.001315    Accuracy:97.702%
Epoch : 4 [56000/60000 (93%)]    Loss: 0.065532    Accuracy:97.717%
Epoch : 4 [57600/60000 (96%)]    Loss: 0.120293    Accuracy:97.729%
Epoch : 4 [59200/60000 (99%)]    Loss: 0.000217    Accuracy:97.758%
```

```
[68]: def evaluate(model):
    #model = mlp
    correct = 0
    for test_imgs, test_labels in test_loader:
        #print(test_imgs.shape)
        test_imgs = Variable(test_imgs).float()
        output = model(test_imgs)
        predicted = torch.max(output,1)[1]
        correct += (predicted == test_labels).sum()
    print("Test accuracy:{:.3f}% ".format( float(correct) /
    ↪(len(test_loader)*BATCH_SIZE)))
    evaluate(mlp)
```

Test accuracy:0.971%

##Since a CNN needs a image shape as input let's reshape our flatten images to real image

```
[69]: torch_X_train = torch_X_train.view(-1, 1,28,28).float()
    torch_X_test = torch_X_test.view(-1,1,28,28).float()
    print(torch_X_train.shape)
    print(torch_X_test.shape)

    # Pytorch train and test sets
    train = torch.utils.data.TensorDataset(torch_X_train,torch_y_train)
    test = torch.utils.data.TensorDataset(torch_X_test,torch_y_test)

    # data loader
    train_loader = torch.utils.data.DataLoader(train, batch_size = BATCH_SIZE,
    ↪shuffle = False)
    test_loader = torch.utils.data.DataLoader(test, batch_size = BATCH_SIZE,
    ↪shuffle = False)MB
```

torch.Size([60000, 1, 28, 28])

torch.Size([10000, 1, 28, 28])

```
[70]: class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=5)
        self.conv2 = nn.Conv2d(32, 32, kernel_size=5)
        self.conv3 = nn.Conv2d(32,64, kernel_size=5)
        self.fc1 = nn.Linear(3*3*64, 256)
```

```

        self.fc2 = nn.Linear(256, 10)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        #x = F.dropout(x, p=0.5, training=self.training)
        x = F.relu(F.max_pool2d(self.conv2(x), 2))
        x = F.dropout(x, p=0.5, training=self.training)
        x = F.relu(F.max_pool2d(self.conv3(x), 2))
        x = F.dropout(x, p=0.5, training=self.training)
        x = x.view(-1, 3*3*64)
        x = F.relu(self.fc1(x))
        x = F.dropout(x, training=self.training)
        x = self.fc2(x)
        return F.log_softmax(x, dim=1)

cnn = CNN()
print(cnn)

it = iter(train_loader)
X_batch, y_batch = next(it)
print(cnn.forward(X_batch).shape)

```

```

CNN(
  (conv1): Conv2d(1, 32, kernel_size=(5, 5), stride=(1, 1))
  (conv2): Conv2d(32, 32, kernel_size=(5, 5), stride=(1, 1))
  (conv3): Conv2d(32, 64, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear(in_features=576, out_features=256, bias=True)
  (fc2): Linear(in_features=256, out_features=10, bias=True)
)
torch.Size([32, 10])

```

```
[71]: fit(cnn, train_loader)
```

Epoch : 0 [0/60000 (0%)]	Loss: 24.703955	Accuracy: 6.250%
Epoch : 0 [1600/60000 (3%)]	Loss: 1.980315	Accuracy: 17.279%
Epoch : 0 [3200/60000 (5%)]	Loss: 1.460487	Accuracy: 29.115%
Epoch : 0 [4800/60000 (8%)]	Loss: 0.748759	Accuracy: 38.949%
Epoch : 0 [6400/60000 (11%)]	Loss: 1.238455	Accuracy: 45.553%
Epoch : 0 [8000/60000 (13%)]	Loss: 0.865389	Accuracy: 50.461%
Epoch : 0 [9600/60000 (16%)]	Loss: 0.764462	Accuracy: 54.485%
Epoch : 0 [11200/60000 (19%)]	Loss: 0.602340	Accuracy: 58.191%
Epoch : 0 [12800/60000 (21%)]	Loss: 0.548422	Accuracy: 60.793%
Epoch : 0 [14400/60000 (24%)]	Loss: 0.491586	Accuracy: 63.089%
Epoch : 0 [16000/60000 (27%)]	Loss: 0.702923	Accuracy: 64.814%
Epoch : 0 [17600/60000 (29%)]	Loss: 0.591373	Accuracy: 66.640%
Epoch : 0 [19200/60000 (32%)]	Loss: 0.427941	Accuracy: 68.225%
Epoch : 0 [20800/60000 (35%)]	Loss: 0.508253	Accuracy: 69.772%
Epoch : 0 [22400/60000 (37%)]	Loss: 0.220711	Accuracy: 71.122%

Epoch : 0	[24000/60000 (40%)]	Loss: 0.182511	Accuracy:72.121%
Epoch : 0	[25600/60000 (43%)]	Loss: 0.578731	Accuracy:73.174%
Epoch : 0	[27200/60000 (45%)]	Loss: 0.405716	Accuracy:74.104%
Epoch : 0	[28800/60000 (48%)]	Loss: 0.335528	Accuracy:75.035%
Epoch : 0	[30400/60000 (51%)]	Loss: 0.524058	Accuracy:75.710%
Epoch : 0	[32000/60000 (53%)]	Loss: 0.522807	Accuracy:76.346%
Epoch : 0	[33600/60000 (56%)]	Loss: 0.353773	Accuracy:77.019%
Epoch : 0	[35200/60000 (59%)]	Loss: 0.348661	Accuracy:77.677%
Epoch : 0	[36800/60000 (61%)]	Loss: 0.335008	Accuracy:78.307%
Epoch : 0	[38400/60000 (64%)]	Loss: 0.295484	Accuracy:78.841%
Epoch : 0	[40000/60000 (67%)]	Loss: 0.330925	Accuracy:79.359%
Epoch : 0	[41600/60000 (69%)]	Loss: 0.761555	Accuracy:79.790%
Epoch : 0	[43200/60000 (72%)]	Loss: 0.308168	Accuracy:80.204%
Epoch : 0	[44800/60000 (75%)]	Loss: 0.039641	Accuracy:80.661%
Epoch : 0	[46400/60000 (77%)]	Loss: 0.347005	Accuracy:81.032%
Epoch : 0	[48000/60000 (80%)]	Loss: 0.695292	Accuracy:81.404%
Epoch : 0	[49600/60000 (83%)]	Loss: 0.664559	Accuracy:81.742%
Epoch : 0	[51200/60000 (85%)]	Loss: 0.267460	Accuracy:82.083%
Epoch : 0	[52800/60000 (88%)]	Loss: 0.457143	Accuracy:82.418%
Epoch : 0	[54400/60000 (91%)]	Loss: 0.123103	Accuracy:82.742%
Epoch : 0	[56000/60000 (93%)]	Loss: 0.171987	Accuracy:83.040%
Epoch : 0	[57600/60000 (96%)]	Loss: 0.293099	Accuracy:83.320%
Epoch : 0	[59200/60000 (99%)]	Loss: 0.139004	Accuracy:83.656%
Epoch : 1	[0/60000 (0%)]	Loss: 0.197437	Accuracy:93.750%
Epoch : 1	[1600/60000 (3%)]	Loss: 0.182057	Accuracy:93.137%
Epoch : 1	[3200/60000 (5%)]	Loss: 0.449420	Accuracy:94.028%
Epoch : 1	[4800/60000 (8%)]	Loss: 0.116307	Accuracy:93.729%
Epoch : 1	[6400/60000 (11%)]	Loss: 0.029951	Accuracy:93.797%
Epoch : 1	[8000/60000 (13%)]	Loss: 0.101050	Accuracy:93.875%
Epoch : 1	[9600/60000 (16%)]	Loss: 0.333615	Accuracy:93.625%
Epoch : 1	[11200/60000 (19%)]	Loss: 0.426801	Accuracy:93.643%
Epoch : 1	[12800/60000 (21%)]	Loss: 0.298322	Accuracy:93.462%
Epoch : 1	[14400/60000 (24%)]	Loss: 0.108020	Accuracy:93.528%
Epoch : 1	[16000/60000 (27%)]	Loss: 0.474114	Accuracy:93.469%
Epoch : 1	[17600/60000 (29%)]	Loss: 0.277469	Accuracy:93.455%
Epoch : 1	[19200/60000 (32%)]	Loss: 0.163615	Accuracy:93.506%
Epoch : 1	[20800/60000 (35%)]	Loss: 0.227395	Accuracy:93.568%
Epoch : 1	[22400/60000 (37%)]	Loss: 0.010190	Accuracy:93.630%
Epoch : 1	[24000/60000 (40%)]	Loss: 0.085058	Accuracy:93.633%
Epoch : 1	[25600/60000 (43%)]	Loss: 0.092341	Accuracy:93.606%
Epoch : 1	[27200/60000 (45%)]	Loss: 0.161767	Accuracy:93.548%
Epoch : 1	[28800/60000 (48%)]	Loss: 0.029576	Accuracy:93.573%
Epoch : 1	[30400/60000 (51%)]	Loss: 0.205116	Accuracy:93.589%
Epoch : 1	[32000/60000 (53%)]	Loss: 0.252389	Accuracy:93.581%
Epoch : 1	[33600/60000 (56%)]	Loss: 0.318719	Accuracy:93.583%
Epoch : 1	[35200/60000 (59%)]	Loss: 0.124187	Accuracy:93.583%
Epoch : 1	[36800/60000 (61%)]	Loss: 0.110199	Accuracy:93.636%
Epoch : 1	[38400/60000 (64%)]	Loss: 0.135114	Accuracy:93.633%

Epoch : 1	[40000/60000 (67%)]	Loss: 0.337480	Accuracy:93.608%
Epoch : 1	[41600/60000 (69%)]	Loss: 0.151904	Accuracy:93.637%
Epoch : 1	[43200/60000 (72%)]	Loss: 0.201784	Accuracy:93.651%
Epoch : 1	[44800/60000 (75%)]	Loss: 0.120788	Accuracy:93.670%
Epoch : 1	[46400/60000 (77%)]	Loss: 0.631979	Accuracy:93.638%
Epoch : 1	[48000/60000 (80%)]	Loss: 0.330198	Accuracy:93.644%
Epoch : 1	[49600/60000 (83%)]	Loss: 0.092436	Accuracy:93.667%
Epoch : 1	[51200/60000 (85%)]	Loss: 0.046334	Accuracy:93.693%
Epoch : 1	[52800/60000 (88%)]	Loss: 0.316572	Accuracy:93.701%
Epoch : 1	[54400/60000 (91%)]	Loss: 0.114109	Accuracy:93.695%
Epoch : 1	[56000/60000 (93%)]	Loss: 0.394695	Accuracy:93.727%
Epoch : 1	[57600/60000 (96%)]	Loss: 0.083526	Accuracy:93.792%
Epoch : 1	[59200/60000 (99%)]	Loss: 0.015199	Accuracy:93.860%
Epoch : 2	[0/60000 (0%)]	Loss: 0.217455	Accuracy:90.625%
Epoch : 2	[1600/60000 (3%)]	Loss: 0.406698	Accuracy:94.056%
Epoch : 2	[3200/60000 (5%)]	Loss: 0.423212	Accuracy:94.864%
Epoch : 2	[4800/60000 (8%)]	Loss: 0.065729	Accuracy:94.454%
Epoch : 2	[6400/60000 (11%)]	Loss: 0.006236	Accuracy:94.325%
Epoch : 2	[8000/60000 (13%)]	Loss: 0.204239	Accuracy:94.410%
Epoch : 2	[9600/60000 (16%)]	Loss: 0.303900	Accuracy:94.248%
Epoch : 2	[11200/60000 (19%)]	Loss: 0.255757	Accuracy:94.436%
Epoch : 2	[12800/60000 (21%)]	Loss: 0.152741	Accuracy:94.459%
Epoch : 2	[14400/60000 (24%)]	Loss: 0.082564	Accuracy:94.471%
Epoch : 2	[16000/60000 (27%)]	Loss: 0.283941	Accuracy:94.374%
Epoch : 2	[17600/60000 (29%)]	Loss: 0.014378	Accuracy:94.380%
Epoch : 2	[19200/60000 (32%)]	Loss: 0.083600	Accuracy:94.452%
Epoch : 2	[20800/60000 (35%)]	Loss: 0.360928	Accuracy:94.556%
Epoch : 2	[22400/60000 (37%)]	Loss: 0.102351	Accuracy:94.593%
Epoch : 2	[24000/60000 (40%)]	Loss: 0.109202	Accuracy:94.653%
Epoch : 2	[25600/60000 (43%)]	Loss: 0.092279	Accuracy:94.714%
Epoch : 2	[27200/60000 (45%)]	Loss: 0.241476	Accuracy:94.679%
Epoch : 2	[28800/60000 (48%)]	Loss: 0.107590	Accuracy:94.704%
Epoch : 2	[30400/60000 (51%)]	Loss: 0.188392	Accuracy:94.759%
Epoch : 2	[32000/60000 (53%)]	Loss: 0.298663	Accuracy:94.699%
Epoch : 2	[33600/60000 (56%)]	Loss: 0.165397	Accuracy:94.690%
Epoch : 2	[35200/60000 (59%)]	Loss: 0.109414	Accuracy:94.738%
Epoch : 2	[36800/60000 (61%)]	Loss: 0.111880	Accuracy:94.790%
Epoch : 2	[38400/60000 (64%)]	Loss: 0.247883	Accuracy:94.786%
Epoch : 2	[40000/60000 (67%)]	Loss: 0.241609	Accuracy:94.799%
Epoch : 2	[41600/60000 (69%)]	Loss: 0.274725	Accuracy:94.809%
Epoch : 2	[43200/60000 (72%)]	Loss: 0.335197	Accuracy:94.828%
Epoch : 2	[44800/60000 (75%)]	Loss: 0.005174	Accuracy:94.872%
Epoch : 2	[46400/60000 (77%)]	Loss: 1.024677	Accuracy:94.889%
Epoch : 2	[48000/60000 (80%)]	Loss: 0.283337	Accuracy:94.860%
Epoch : 2	[49600/60000 (83%)]	Loss: 0.264172	Accuracy:94.854%
Epoch : 2	[51200/60000 (85%)]	Loss: 0.145659	Accuracy:94.857%
Epoch : 2	[52800/60000 (88%)]	Loss: 0.662844	Accuracy:94.893%
Epoch : 2	[54400/60000 (91%)]	Loss: 0.017660	Accuracy:94.922%

Epoch : 2	[56000/60000 (93%)]	Loss: 0.118975	Accuracy:94.960%
Epoch : 2	[57600/60000 (96%)]	Loss: 0.390172	Accuracy:94.961%
Epoch : 2	[59200/60000 (99%)]	Loss: 0.000941	Accuracy:95.016%
Epoch : 3	[0/60000 (0%)]	Loss: 0.039304	Accuracy:96.875%
Epoch : 3	[1600/60000 (3%)]	Loss: 0.274808	Accuracy:94.424%
Epoch : 3	[3200/60000 (5%)]	Loss: 0.082404	Accuracy:95.142%
Epoch : 3	[4800/60000 (8%)]	Loss: 0.180459	Accuracy:95.281%
Epoch : 3	[6400/60000 (11%)]	Loss: 0.055648	Accuracy:95.134%
Epoch : 3	[8000/60000 (13%)]	Loss: 0.032606	Accuracy:95.219%
Epoch : 3	[9600/60000 (16%)]	Loss: 0.057986	Accuracy:94.975%
Epoch : 3	[11200/60000 (19%)]	Loss: 0.275520	Accuracy:95.103%
Epoch : 3	[12800/60000 (21%)]	Loss: 0.070716	Accuracy:94.950%
Epoch : 3	[14400/60000 (24%)]	Loss: 0.027030	Accuracy:94.928%
Epoch : 3	[16000/60000 (27%)]	Loss: 0.066866	Accuracy:94.941%
Epoch : 3	[17600/60000 (29%)]	Loss: 0.028506	Accuracy:94.941%
Epoch : 3	[19200/60000 (32%)]	Loss: 0.219603	Accuracy:94.998%
Epoch : 3	[20800/60000 (35%)]	Loss: 0.051435	Accuracy:94.993%
Epoch : 3	[22400/60000 (37%)]	Loss: 0.011864	Accuracy:95.047%
Epoch : 3	[24000/60000 (40%)]	Loss: 0.177396	Accuracy:95.073%
Epoch : 3	[25600/60000 (43%)]	Loss: 0.050125	Accuracy:95.154%
Epoch : 3	[27200/60000 (45%)]	Loss: 0.215653	Accuracy:95.131%
Epoch : 3	[28800/60000 (48%)]	Loss: 0.357761	Accuracy:95.106%
Epoch : 3	[30400/60000 (51%)]	Loss: 0.303260	Accuracy:95.127%
Epoch : 3	[32000/60000 (53%)]	Loss: 0.435719	Accuracy:95.042%
Epoch : 3	[33600/60000 (56%)]	Loss: 0.038899	Accuracy:95.008%
Epoch : 3	[35200/60000 (59%)]	Loss: 0.157955	Accuracy:95.084%
Epoch : 3	[36800/60000 (61%)]	Loss: 0.065332	Accuracy:95.127%
Epoch : 3	[38400/60000 (64%)]	Loss: 0.087521	Accuracy:95.145%
Epoch : 3	[40000/60000 (67%)]	Loss: 0.119639	Accuracy:95.131%
Epoch : 3	[41600/60000 (69%)]	Loss: 0.089203	Accuracy:95.143%
Epoch : 3	[43200/60000 (72%)]	Loss: 0.412446	Accuracy:95.152%
Epoch : 3	[44800/60000 (75%)]	Loss: 0.103714	Accuracy:95.211%
Epoch : 3	[46400/60000 (77%)]	Loss: 0.355738	Accuracy:95.202%
Epoch : 3	[48000/60000 (80%)]	Loss: 0.139295	Accuracy:95.191%
Epoch : 3	[49600/60000 (83%)]	Loss: 0.194295	Accuracy:95.189%
Epoch : 3	[51200/60000 (85%)]	Loss: 0.083338	Accuracy:95.181%
Epoch : 3	[52800/60000 (88%)]	Loss: 0.779976	Accuracy:95.192%
Epoch : 3	[54400/60000 (91%)]	Loss: 0.033422	Accuracy:95.196%
Epoch : 3	[56000/60000 (93%)]	Loss: 0.117889	Accuracy:95.208%
Epoch : 3	[57600/60000 (96%)]	Loss: 0.146757	Accuracy:95.232%
Epoch : 3	[59200/60000 (99%)]	Loss: 0.007108	Accuracy:95.286%
Epoch : 4	[0/60000 (0%)]	Loss: 0.037586	Accuracy:96.875%
Epoch : 4	[1600/60000 (3%)]	Loss: 0.532983	Accuracy:95.282%
Epoch : 4	[3200/60000 (5%)]	Loss: 0.141879	Accuracy:95.575%
Epoch : 4	[4800/60000 (8%)]	Loss: 0.068667	Accuracy:95.882%
Epoch : 4	[6400/60000 (11%)]	Loss: 0.088112	Accuracy:95.740%
Epoch : 4	[8000/60000 (13%)]	Loss: 0.112370	Accuracy:95.754%
Epoch : 4	[9600/60000 (16%)]	Loss: 0.098440	Accuracy:95.650%

Epoch : 4	[11200/60000 (19%)]	Loss: 0.284854	Accuracy:95.593%
Epoch : 4	[12800/60000 (21%)]	Loss: 0.082128	Accuracy:95.605%
Epoch : 4	[14400/60000 (24%)]	Loss: 0.066085	Accuracy:95.468%
Epoch : 4	[16000/60000 (27%)]	Loss: 0.137398	Accuracy:95.422%
Epoch : 4	[17600/60000 (29%)]	Loss: 0.083179	Accuracy:95.389%
Epoch : 4	[19200/60000 (32%)]	Loss: 0.296649	Accuracy:95.372%
Epoch : 4	[20800/60000 (35%)]	Loss: 0.079538	Accuracy:95.387%
Epoch : 4	[22400/60000 (37%)]	Loss: 0.021219	Accuracy:95.399%
Epoch : 4	[24000/60000 (40%)]	Loss: 0.045635	Accuracy:95.460%
Epoch : 4	[25600/60000 (43%)]	Loss: 0.025185	Accuracy:95.431%
Epoch : 4	[27200/60000 (45%)]	Loss: 0.313549	Accuracy:95.384%
Epoch : 4	[28800/60000 (48%)]	Loss: 0.358900	Accuracy:95.384%
Epoch : 4	[30400/60000 (51%)]	Loss: 0.290802	Accuracy:95.354%
Epoch : 4	[32000/60000 (53%)]	Loss: 0.031017	Accuracy:95.367%
Epoch : 4	[33600/60000 (56%)]	Loss: 0.085758	Accuracy:95.362%
Epoch : 4	[35200/60000 (59%)]	Loss: 0.152539	Accuracy:95.416%
Epoch : 4	[36800/60000 (61%)]	Loss: 0.162986	Accuracy:95.441%
Epoch : 4	[38400/60000 (64%)]	Loss: 0.174990	Accuracy:95.444%
Epoch : 4	[40000/60000 (67%)]	Loss: 0.521610	Accuracy:95.456%
Epoch : 4	[41600/60000 (69%)]	Loss: 0.122210	Accuracy:95.484%
Epoch : 4	[43200/60000 (72%)]	Loss: 0.318961	Accuracy:95.503%
Epoch : 4	[44800/60000 (75%)]	Loss: 0.035795	Accuracy:95.550%
Epoch : 4	[46400/60000 (77%)]	Loss: 0.384703	Accuracy:95.535%
Epoch : 4	[48000/60000 (80%)]	Loss: 0.491851	Accuracy:95.561%
Epoch : 4	[49600/60000 (83%)]	Loss: 0.055732	Accuracy:95.561%
Epoch : 4	[51200/60000 (85%)]	Loss: 0.087599	Accuracy:95.583%
Epoch : 4	[52800/60000 (88%)]	Loss: 0.285666	Accuracy:95.597%
Epoch : 4	[54400/60000 (91%)]	Loss: 0.043655	Accuracy:95.602%
Epoch : 4	[56000/60000 (93%)]	Loss: 0.363469	Accuracy:95.604%
Epoch : 4	[57600/60000 (96%)]	Loss: 0.103343	Accuracy:95.624%
Epoch : 4	[59200/60000 (99%)]	Loss: 0.018067	Accuracy:95.676%

0.3 SVM Classifier for MNIST DATA SET

```
[0]: from keras.datasets import mnist
      (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
[0]: x_train = x_train.reshape(60000,784)
      x_test = x_test.reshape(10000,784)
```

```
[0]: ## Applying HOG feature extraction
```

```
[0]: from sklearn.svm import SVC
      from sklearn.metrics import accuracy_score
      from skimage.feature import hog
      from sklearn import preprocessing
```

```
from collections import Counter
```

```
[38]: list_hog_train = []  
for feature in x_train:  
    fd = hog(feature.reshape((28,28)), orientations=10,  
    ↪pixels_per_cell=(7,7),cells_per_block=(1,1),visualize=False )  
    list_hog_train.append(fd)  
hog_features = np.array(list_hog_train, 'float64')  
preProcess = preprocessing.MaxAbsScaler().fit(hog_features)  
hog_features_transformed_train = preProcess.transform(hog_features)  
print(hog_features_transformed_train.shape)
```

(60000, 160)

```
[0]: ## Extracting hog feature for test data
```

```
[41]: list_hog_test = []  
for feature in x_test:  
    fd = hog(feature.reshape((28,28)), orientations=10,  
    ↪pixels_per_cell=(7,7),cells_per_block=(1,1),visualize=False )  
    list_hog_test.append(fd)  
hog_features_test = np.array(list_hog_test, 'float64')  
preProcess = preprocessing.MaxAbsScaler().fit(hog_features_test)  
hog_features_transformed_test = preProcess.transform(hog_features_test)  
print(hog_features_transformed_test.shape)
```

(10000, 160)

```
[0]: model = SVC()  
model.fit(hog_features_transformed_train,y_train)  
y_pred = model.predict(hog_features_transformed_test)
```

Fitting the model with best parameters

```
[43]: print(accuracy_score(y_test, y_pred))
```

0.9713

q4

April 3, 2020

1 Question-4 House Electricity Consumption Prediction

```
[0]: from zipfile import ZipFile
file_name="household_power_consumption.zip"

with ZipFile(file_name,'r') as zip:
    zip.extractall()
    print('Done')

# !unzip /content/household_power_consumption.zip
```

Done

2 Reading the data file

```
[0]: %tensorflow_version 1.x
import pandas as pd

headers = ['Date', 'Time', 'Global_active_power', 'Global_reactive_power',
           'Voltage', 'Global_intensity', 'Sub_metering_1', 'Sub_metering_2',
           'Sub_metering_3']

dtypes = {'Date':'str', 'Time':'str', 'Global_active_power':'float',
          'Global_reactive_power': 'float', 'Voltage':'float',
          'Global_intensity':'float', 'Sub_metering_1':'float',
          'Sub_metering_2':'float', 'Sub_metering_3':'float'}

df = pd.read_csv('household_power_consumption.txt', sep=';',
                 dtype=dtypes, na_values=['?'])
df.head()
```

```
[0]:
```

	Date	Time	...	Sub_metering_2	Sub_metering_3
0	16/12/2006	17:24:00	...	1.0	17.0
1	16/12/2006	17:25:00	...	1.0	16.0
2	16/12/2006	17:26:00	...	2.0	17.0

```

3  16/12/2006  17:27:00  ...          1.0          17.0
4  16/12/2006  17:28:00  ...          1.0          17.0

```

[5 rows x 9 columns]

```

[0]: import numpy as np
      from multiprocessing import cpu_count, Pool

      def parallel_map(data, func):
          n_cores = cpu_count()
          data_split = np.array_split(data, n_cores)
          pool = Pool(n_cores)
          data = pd.concat(pool.map(func, data_split))
          pool.close()
          pool.join()
          return data

      def parse(row):
          row['DateTime'] = pd.to_datetime(row['DateTime'],
                                           format='%d/%m/%Y %H:%M:%S')

          return row

```

```

[0]: df['DateTime'] = df['Date'] + ' ' + df['Time']
      df = parallel_map(df, parse)
      df.dtypes

```

```

[0]: Date          object
      Time          object
      Global_active_power    float64
      Global_reactive_power  float64
      Voltage                float64
      Global_intensity        float64
      Sub_metering_1          float64
      Sub_metering_2          float64
      Sub_metering_3          float64
      DateTime              datetime64[ns]
      dtype: object

```

```

[0]: df.drop(['Date', 'Time'], axis=1, inplace=True)
      df = df[[df.columns[-1]] + list(df.columns[:-1])]
      df.set_index('DateTime', inplace=True)
      df.head()

```

```

[0]:          Global_active_power  ...  Sub_metering_3
      DateTime
2006-12-16 17:24:00          4.216  ...          17.0
2006-12-16 17:25:00          5.360  ...          16.0

```

2006-12-16 17:26:00	5.374	...	17.0
2006-12-16 17:27:00	5.388	...	17.0
2006-12-16 17:28:00	3.666	...	17.0

[5 rows x 7 columns]


```
[0]: df['hour'] = df.index.hour
df['day'] = df.index.day
df['month'] = df.index.month
df['day_of_week'] = df.index.dayofweek
```

```
[0]: df['Rest_active_power'] = df['Global_active_power'] * 1000 / 60 - \
df['Sub_metering_1'] - df['Sub_metering_2'] - \
df['Sub_metering_3']
```

3 Identifying and handling Missing Data


```
[0]: import seaborn as sns
sns.distplot(df)
```

```
[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7fbcef5a8358>
```



```
[0]: sns.heatmap(df.isnull(),yticklabels=False,cbar=False)
```

```
[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7fbd67132198>
```



```
[0]: number = len(df) - len(df.dropna())
percentage = number * 100 / len(df)
print(f'Number of points with missing values is: {number}')
print(f'Percentage of points with missing values is: {percentage}\n')
print(f'Missing value counts:\n{df.isnull().sum(axis=0)}\n')
```

Number of points with missing values is: 25979

Percentage of points with missing values is: 1.2518437457686005

Missing value counts:

Global_active_power	25979
Global_reactive_power	25979
Voltage	25979
Global_intensity	25979
Sub_metering_1	25979
Sub_metering_2	25979

```

Sub_metering_3      25979
hour                 0
day                  0
month                0
day_of_week          0
Rest_active_power    25979
dtype: int64

```

#Power consumption over the whole timespan

Average global active power over resampled data yearly, quarterly, monthly, and daily.

```

[0]: # Tableau colors
tableau20 = [(31, 119, 180), (174, 199, 232), (255, 127, 14), (255, 187, 120),
             (44, 160, 44), (152, 223, 138), (214, 39, 40), (255, 152, 150),
             (148, 103, 189), (197, 176, 213), (140, 86, 75), (196, 156, 148),
             (227, 119, 194), (247, 182, 210), (127, 127, 127), (199, 199, 199),
             (188, 189, 34), (219, 219, 141), (23, 190, 207), (158, 218, 229)]

# Scale the RGB values to the [0, 1] range, which is the format matplotlib
→ accepts.
for i in range(len(tableau20)):
    r, g, b = tableau20[i]
    tableau20[i] = (r / 255., g / 255., b / 255.)

```

```

[0]: import matplotlib.pyplot as plt
import matplotlib.dates as mdates
from pandas.plotting import register_matplotlib_converters
register_matplotlib_converters()
my_fmt = mdates.DateFormatter('%a %d/%m %H:%M')

fig, ax = plt.subplots(nrows=2, ncols=2, figsize=(22, 22))

frequencies = ['1Y', '1M', '1Q', '1D']

dic = {'1Y': 'Yearly', '1M': 'Monthly', '1Q': 'Quarterly', '1D': 'Daily'}


i = 0
for row in ax:
    for col in row:
        for tick in col.get_xticklabels():
            tick.set_rotation(45)
        col.plot(df[['Global_active_power']].resample(frequencies[i]).mean(),
                 color=tableau20[i * 2])
        col.set_xlabel('Time')
        col.set_ylabel('Global Active Power (kW)')
        col.set_title(dic[frequencies[i]])

```



```
col.xaxis.set_major_formatter(my_fmt)

# Aesthetics
col.spines["top"].set_visible(False)
col.spines["right"].set_visible(False)
i += 1
```



4 supervised learning

Predict global active power at time $t + h$ given all the variables at times $t, t - 1, \dots, t - w$, where h is the prediction horizon and w is the window size. Here's a function that takes as input a time

series dataframe, a window size, a horizon, a set of variables to be lagged and a set of variables to be forecasted and outputs the dataframe transformed and ready for supervised learning.

```
[0]: from pandas import DataFrame
from pandas import concat

def series_to_supervised(data, window_size=1, horizon=1, inputs='all',
    ↪targets='all'):
    """
    Frame a time series as a supervised learning dataset.

    Arguments:
        data: A pandas DataFrame containing the time series
              (the index must be a DateTimeIndex).
        window_size: Number of lagged observations as input.
        horizon: Number of steps to forecast ahead.
        inputs: A list of the columns of the dataframe to be lagged.
        targets: A list of the columns of the dataframe to be forecasted.

    Returns:
        Pandas DataFrame of series framed for supervised learning.
    """

    if targets == 'all':
        targets = data.columns

    if inputs == 'all':
        inputs = data.columns

    result = DataFrame(index=df.index)
    names = []

    # input sequence (t-w, ..., t-1)
    for i in range(window_size, 0, -1):
        result = pd.concat([result, data[inputs].shift(i)], axis=1)
        names += [(f'{data[inputs].columns[j]}(t-{i})') for j in
    ↪range(len(inputs))]

    # the input not shifted (t)
    result = pd.concat([result, data.copy()], axis=1)
    names += [(f'{column}(t)') for column in data.columns]

    # forecast (t+h)
    for i in [horizon]:
        result = pd.concat([result, data[targets].shift(-i)], axis=1)
```

```

        names += [(f'{data[targets].columns[j]}(t+{i})') for j in
↪range(len(targets))]

    # put it all together
    result.columns = names

    # drop rows with NaN values
    result.dropna(inplace=True)
    return result

```

```

[0]: inputs = ['Global_active_power', 'Global_reactive_power', 'Voltage',
              'Global_intensity', 'Sub_metering_1', 'Sub_metering_2',
              'Sub_metering_3', 'Rest_active_power']

targets = ['Global_active_power']

df_supervised = series_to_supervised(df, window_size=5, horizon=1,
↪inputs=inputs, targets=targets)
df_supervised.head()

```

```

[0]:
           Global_active_power(t-5)  ...  Global_active_power(t+1)
DateTime
2006-12-16 17:29:00                4.216  ...                3.702
2006-12-16 17:30:00                5.360  ...                3.700
2006-12-16 17:31:00                5.374  ...                3.668
2006-12-16 17:32:00                5.388  ...                3.662
2006-12-16 17:33:00                3.666  ...                4.448

[5 rows x 53 columns]

```

```

[0]: #storing into some file to visualize at different window size
      #df_supervised.to_csv('supervised_w10_h1.csv', index=True)

```

5 Supervised Machine Learning

```

[0]: # import pandas as pd

      # df_supervised = pd.read_csv('supervised_w10_h1.csv', parse_dates=['DateTime'])
      # df_supervised.set_index('DateTime', inplace=True)

```

6 Splitting data

splitting data into train,validate and test

```
[0]: def train_validate_test_split(df, train_percent=.6, validate_percent=.2, seed=None):  
    np.random.seed(seed)  
    m = len(df)  
    train_end = int(train_percent * m)  
    validate_end = int(validate_percent * m) + train_end  
    train = df.iloc[:train_end]  
    validate = df.iloc[train_end:validate_end]  
    test = df.iloc[validate_end:]  
    return train, validate, test
```

```
[0]: train, validate, test = train_validate_test_split(df_supervised)
```

```
[0]: print(type(train))  
train.shape
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
[0]: (1229311, 53)
```

```
[0]: X_train = train.values[:, :-1]  
y_train = train.values[:, -1]  
  
X_validate = validate.values[:, :-1]  
y_validate = validate.values[:, -1]  
  
X_test = test.values[:, :-1]  
y_test = test.values[:, -1]
```

```
[0]: X_test.shape
```

```
[0]: (409772, 52)
```

```
[0]: X_train.shape
```

```
[0]: (1229311, 52)
```

```
[0]: X_validate.shape
```

```
[0]: (409770, 52)
```

7 Mean absolute percentage Error

```
[0]: # #from sklearn.utils import check_arrays
# def mean_absolute_percentage_error(y_true, y_pred):
#     #y_true, y_pred = check_arrays(y_true, y_pred)
#     return np.mean(np.abs((y_true - y_pred) / y_true)) * 100
```

```
[0]: from sklearn.metrics import mean_squared_error
from math import sqrt
from sklearn.metrics import accuracy_score
```

```
[0]: import numpy as np

def mean_absolute_percentage_error(y_true, y_pred):
    y_true, y_pred = np.array(y_true), np.array(y_pred)
    return np.mean(np.abs((y_true - y_pred) / y_true)) * 100
```

8 Linear Regression

```
[0]: from sklearn.linear_model import LinearRegression

model = LinearRegression()
model.fit(X_train, y_train)
predictions = model.predict(X_validate)
```

```
[0]: from sklearn.metrics import mean_squared_error

mean_squared_error(predictions, y_validate)
```

```
[0]: 0.055878991130430315
```

```
[0]: mean_absolute_percentage_error(y_validate, predictions)
```

```
[0]: 10.996608631396205
```

```
[0]: rms = sqrt(mean_squared_error(y_validate, predictions))
print(rms)
```

```
0.23638737515026118
```

```
[0]: #accuracy_score(y_validate, predictions, normalize=False)
```

9 Multilayer Perceptron

```
[0]: from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
X_train = scaler.fit_transform(X_train)
```

```
X_validate = scaler.transform(X_validate)
```

```
X_test = scaler.transform(X_test)
```

```
[0]: from keras.layers import Input, Dense, Dropout, LSTM, Reshape, Flatten
```

```
from keras import Sequential
```

```
#from tensorflow.keras.optimizers import SGD
```

```
from keras.callbacks import EarlyStopping
```

```
model = Sequential()
```

```
model.add(Dense(100, activation='relu', input_shape=(X_train.shape[1], )))
```

```
model.add(Dropout(0.2))
```

```
model.add(Dense(1))
```

Using TensorFlow backend.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:66: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4432: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:148: The name tf.placeholder_with_default is deprecated. Please use tf.compat.v1.placeholder_with_default instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3733: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

```
[0]: from keras.optimizers import Adam
model.compile(loss='mean_squared_error',
              optimizer=Adam(lr=0.001))

history = model.fit(X_train, y_train, batch_size=1024, epochs=100,
                  verbose=1,
                  validation_data=(X_validate, y_validate),
                  callbacks=[EarlyStopping(patience=1)])
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1033: The name tf.assign_add is deprecated. Please use tf.compat.v1.assign_add instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1020: The name tf.assign is deprecated. Please use tf.compat.v1.assign instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3005: The name tf.Session is deprecated. Please use tf.compat.v1.Session instead.

Train on 1229311 samples, validate on 409770 samples
Epoch 1/100

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:190: The name tf.get_default_session is deprecated. Please use tf.compat.v1.get_default_session instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:197: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:207: The name tf.global_variables is deprecated. Please use tf.compat.v1.global_variables instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:216: The name tf.is_variable_initialized is deprecated. Please use tf.compat.v1.is_variable_initialized instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:223: The name

tf.variables_initializer is deprecated. Please use
tf.compat.v1.variables_initializer instead.

```
1229311/1229311 [=====] - 7s 6us/step - loss: 0.1475 -  
val_loss: 0.0584  
Epoch 2/100  
1229311/1229311 [=====] - 7s 5us/step - loss: 0.0833 -  
val_loss: 0.0542  
Epoch 3/100  
1229311/1229311 [=====] - 6s 5us/step - loss: 0.0781 -  
val_loss: 0.0521  
Epoch 4/100  
1229311/1229311 [=====] - 7s 5us/step - loss: 0.0756 -  
val_loss: 0.0515  
Epoch 5/100  
1229311/1229311 [=====] - 7s 5us/step - loss: 0.0743 -  
val_loss: 0.0506  
Epoch 6/100  
1229311/1229311 [=====] - 7s 5us/step - loss: 0.0737 -  
val_loss: 0.0506  
Epoch 7/100  
1229311/1229311 [=====] - 7s 5us/step - loss: 0.0731 -  
val_loss: 0.0508
```

```
[0]: predictions = model.predict(X_validate)
```

```
[0]: mean_squared_error(predictions, y_validate)
```

```
[0]: 0.05082758861318869
```

```
[0]: #mean_absolute_percentage_error(y_validate, predictions)
```

```
[0]: rms = sqrt(mean_squared_error(y_validate, predictions))  
print(rms)
```

```
0.2254497474232089
```

10 Visualizing the data

Predicting for the test data

```
[0]: predictions = model.predict(X_test)
```

```
[0]: df_to_plot = test[['Global_active_power(t+1)']].copy()  
df_to_plot['Global_active_power(t+1)_predicted'] = predictions
```



```
[0]: import matplotlib.pyplot as plt
```

```
df_to_plot[:1000].plot()  
plt.show()
```

